

COMP5349 – Cloud Computing

Week 7: Data Flow Engines for Cloud Analytics

A/Prof Dr Uwe Röhm
School of Information Technologies



Lecture Outlook

■ Today:

- ▶ Help with Assignment 1 in the lab rooms in SIT
- ▶ Submission by Friday (tomorrow) 6pm in Blackboard
- ▶ Please also fill in the **Self-Reflection Survey** **after** finishing the assignment

■ Next Week: Details on Spark and Flink

- ▶ Including lab
- ▶ A2 to be published in Week 8

■ Friday this week (28thApril): Data Centre Excursion

- ▶ Visit to one of the Equinix data centres in Alexandria (200 Bourke Rd)
 - ▶ 2 tours @ 1-1:30pm and @ 1:45pm-2:30pm
- => ENROL IN TOUR GROUPS IN Blackboard (under 'Groups')**

Equinix Field Trip

- Tomorrow, Friday 28th April: Data Centre Excursion
 - ▶ 2 tours @ 1-1:30pm and @ 1:45pm-2:30pm
 - ▶ First-Come-First-Served
 - ▶ **ENROL to TOUR GROUPS IN Blackboard ('Groups' section)**
- Requirements
 - ▶ 18+ and an **official Australian photo ID** (drivers license; passport)
 - International students: need to bring passport
 - Note: your student ID card is not accepted
 - ▶ Closed-in shoes
 - ▶ Signed Fieldtrip Acknowledgement Form
- Location: 200 Bourke Road, Alexandria
 - ▶ Public Trapo: 10 minutes walk from Mascot Train Station
 - ▶ Car... Limited parking at Bunnings which is next door



Outline

- Motivation
- Overview of Apache Spark
- Overview of Apache Flink
- Conclusions

COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1969

WARNING

This material has been reproduced and communicated to you by or on behalf of the **University of Sydney** pursuant to Part VB of the Copyright Act 1968 (the Act).

The material in this communication may be subject to copyright under the Act. Any further reproduction or communication of this material by you may be the subject of copyright protection under the Act.

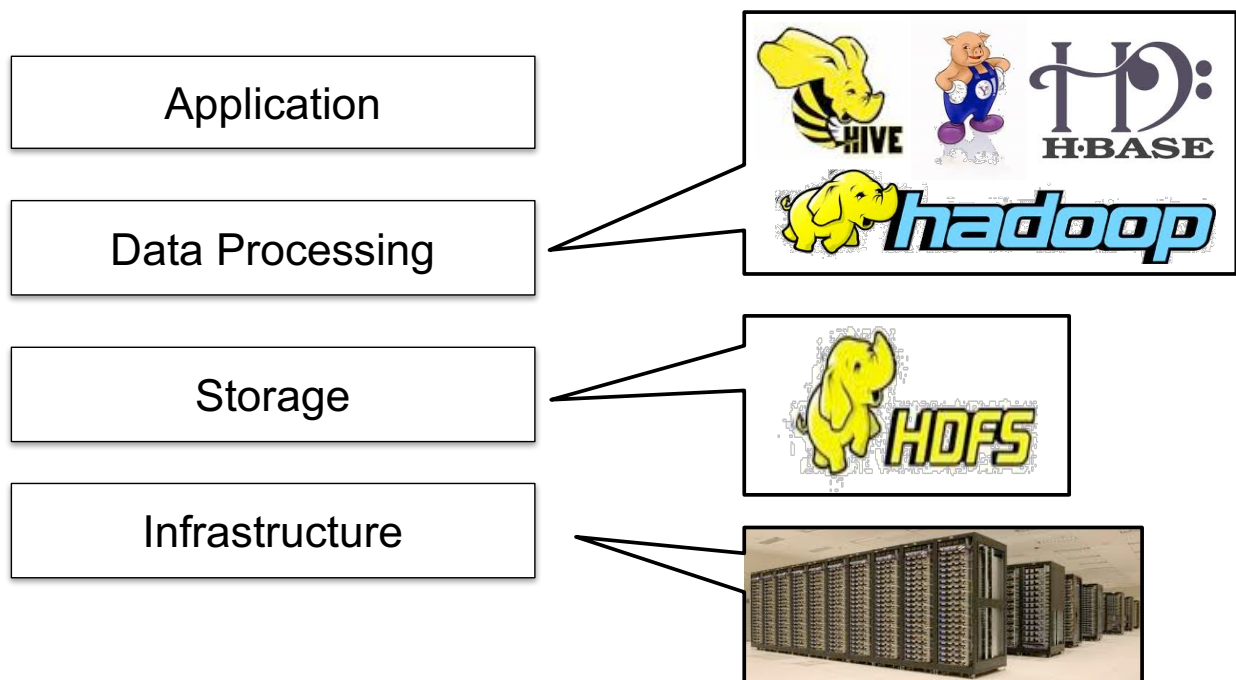
Do not remove this notice

Based on slide decks from Ion Stoica, Mike Franklin,
as well as the online documentation from BDAS, Stratosphere and Flink.




Original Cloud Analytics Stack

- ... mostly focused on large, on-disk datasets: great for **batch**, but **slow**




Motivation

- Map/Reduce allows simple parallelization of pipelined tasks
 - Pro: parallelism, fault-tolerance, runtime can decide where to run tasks
 - Con: simple processing model that assumes data flowing from stable storage to stable storage + materialization of intermediate results



- Can't we do better?




Hive / Pig?

- There are several approaches to provide SQL-like query languages on top of map reduce
- Pro:
 - ▶ High-level abstraction of execution
 - ▶ Data centric view
 - ▶ SQL
- Con:
 - ▶ Core execution model still Map/Reduce jobs
 - ▶ Add significant execution overhead
 - ▶ Expressiveness of SQL is limited; e.g. no machine learning or graph processing possible



Goals



- How to combine **batch**, **streaming**, and **interactive** computations?
- Need for support for more sophisticated algorithms, and easy of use.
- Ideally: compatible with existing infrastructure (Hadoop/HDFS)



Approach

- Extensive use of main **memory** for processing
 - ▶ Rationale: memory is cheap nowadays and much faster than disk (even than SSDs)
 - ▶ Many datasets already fit into memory
 - The inputs of over 90% of jobs in Facebook, Yahoo!, and Bing clusters fit into memory [cite]
- Acyclic data flow plan with advanced operators
 - ▶ Go beyond simple map/reduce for more expressive tasks
- Increased parallelism
 - ▶ Ideally via automated plan optimisation and scheduling
- Built on top of Hadoop/HDFS
 - ▶ Usable with existing jobs and data stores



Berkeley Data Analytics Stack (BDAS)

<https://amplab.cs.berkeley.edu/bdas/>

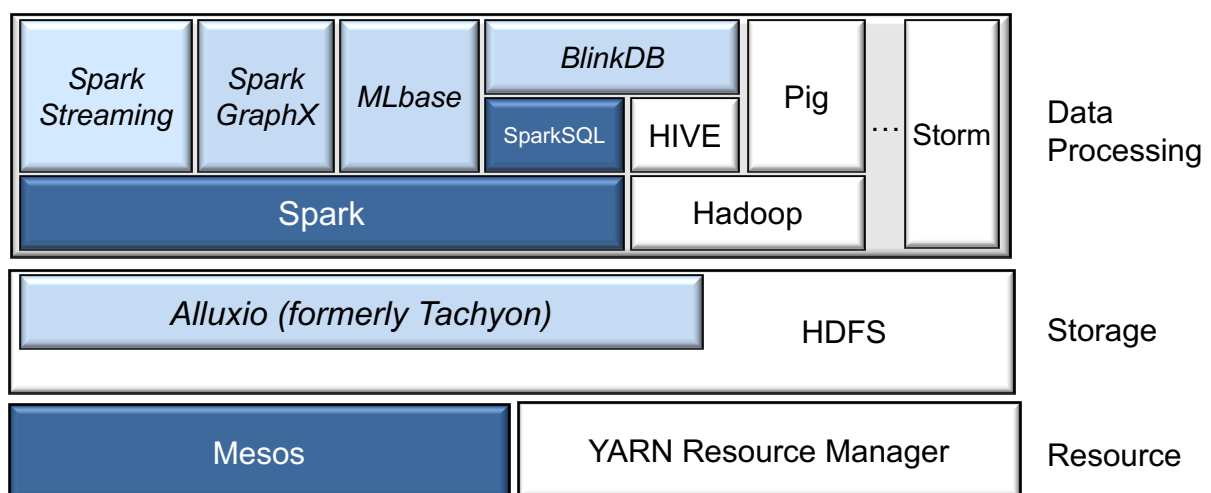


BDAS Approach

- Project by Amplab, UC Berkeley
- Single execution model that supports *batch*, *streaming*, and *interactive* computation models
- Easy to develop sophisticated algorithms
 - ▶ Powerful Python and Scala shells
 - ▶ High level abstractions for graph based, and ML algorithms
- Compatible with existing open source ecosystem (Hadoop/HDFS)
 - ▶ Interoperate with existing storage and input formats (e.g., HDFS, Hive, Flume, ..)
 - ▶ Support existing execution models (e.g., Hive, GraphLab)



Berkeley Data Analytics Stack



- Several extensions based on Hadoop/HDFS
- **Mesos:** Multi-tenant Resource Management
 - ▶ Management platform that allows multiple framework to share cluster; used by eg. Twitter and airbnb
- **Spark:** In-memory framework for interactive and iterative computations
 - ▶ Scala interface, Java and Python APIs
- **SparkSQL:** "HIVE over Spark" - a SQL-like interface, compatible to HIVE queries
- Many more in alpha: Alluxio (in-memory storage), Streaming, MLBase, BlinkDB, ...



Apache Spark

- In-memory framework for interactive and iterative computations
- Goals:
 - ▶ distributed memory abstractions for clusters to support apps with working sets
 - ▶ Retain the attractive properties of MapReduce:
 - Fault tolerance (for crashes & stragglers)
 - Data locality
 - Scalability
- Approach:
 - ▶ augment data flow model with **Resilient Distributed Dataset (RDD)**
 - RDD: fault-tolerance, in-memory storage abstraction



Spark Programming Model

- **Resilient distributed datasets (RDDs)**
 - ▶ Immutable collections partitioned across cluster that can be rebuilt if a partition is lost
 - ▶ Created by transforming data in stable storage using data flow operators (map, filter, group-by, ...)
 - ▶ Can be *cached* across parallel operations
- Parallel operations on RDDs
 - ▶ Reduce, collect, count, save, ...
- Restricted shared variables
 - ▶ Accumulators, broadcast variables



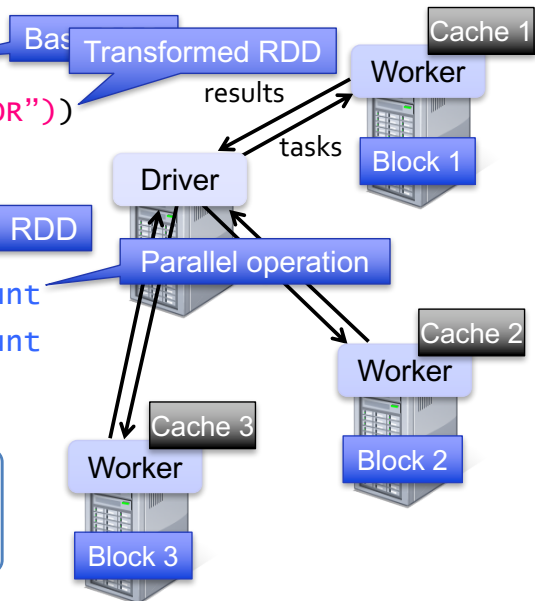
Example: Log Mining

- Load error messages from a log into memory, then interactively search for various patterns

```
lines = spark.textFile("hdfs://...")
errors = lines.filter(_.startsWith("ERROR"))
messages = errors.map(_.split('\t')(2))
cachedMsgs = messages.cache()

cachedMsgs.filter(_.contains("foo")).count
cachedMsgs.filter(_.contains("bar")).count
. . .
```

Result: full-text search of Wikipedia in <1 sec (vs 20 sec for on-disk data)



RDDs in More Detail

- RDD: immutable (read-only), partitioned, logical collection of records
 - Need not be materialized, but rather contains information to rebuild a dataset from stable storage
- Partitioning can be based on a key in each record
 - hash or range partitioning
- Built using bulk transformations on other RDDs
- Can be cached for future reuse
- Operations:

Transformations (define a new RDD)
map
filter
sample
union
groupByKey
reduceByKey
join
cache
...

Parallel operations (return a result to driver)
reduce
collect
count
save
lookupKey
...




RDD Fault Tolerance

- RDDs maintain **lineage** information that can be used to reconstruct lost partitions

- Example:

```
cachedMsgs = textFile(...).filter(_.contains("error"))  
                                .map(_.split('\t')(2))  
                                .cache()
```




Benefits of RDD Model

- Consistency is easy due to immutability
- Inexpensive fault tolerance (log lineage rather than replicating/checkpointing data)
- Locality-aware scheduling of tasks on partitions
- Despite being restricted, model seems applicable to a broad variety of applications



Example App: Logistic Regression

- Goal: find best line separating two sets of points



Example Logistic Regression Code

Load data in memory once

```
val data = spark.textFile(...).map(readPoint).cache()
```

```
var w = Vector.random(D)
```

Initial parameter vector


```
for (i <- 1 to ITERATIONS) {  
  val gradient = data.map(p =>  
    (1 / (1 + exp(-p.y*(w dot p.x))) - 1) * p.y * p.x  
  ).reduce(_ + _)  
  w -= gradient  
}
```

Repeated MapReduce steps
to do gradient descent

```
println("Final w: " + w)
```



Logistic Regression Performance



29 GB dataset on 20 EC2 m1.xlarge machines (4 cores each)



Example: MapReduce

- MapReduce data flow can be expressed using RDD transformations

```
res = data.flatMap(rec => myMapFunc(rec))
            .groupByKey()
            .map((key, vals) => myReduceFunc(key, vals))
```

Or with combiners:

```
res = data.flatMap(rec => myMapFunc(rec))
            .reduceByKey(myCombiner)
            .map((key, val) => myReduceFunc(key, val))
```



Word Count in Spark

```
val lines = spark.textFile("hdfs://...")

val counts = lines.flatMap(_.split("\\s"))
                    .reduceByKey(_ + _)

counts.save("hdfs://...")
```



BDAS: Work in Progress

- **Tachyon** [alpha release so far]
 - ▶ High-throughput, fault-tolerant in-memory storage
 - ▶ Interface compatible to HDFS; support for Spark and Hadoop
- **BlinkDB** [alpha release]
 - ▶ large scale approximate query engine
 - ▶ Allow users to specify error or time bounds
- **SparkGraph** [alpha release]
 - ▶ GraphLab API and Toolkits on top of Spark
 - ▶ Fault tolerance by leveraging Spark
- **MLbase** [in development]
 - ▶ Declarative approach to ML
 - ▶ Develop scalable ML algorithms
 - ▶ Make ML accessible to non-experts



Apache Flink

<http://flink.apache.org>



Apache Flink Project

- Started as a research project ('Stratosphere') by database research groups at TU Berlin, Humboldt University Berlin and HPI Potsdam
- Since Feb 2015 Apache top-level project (Apache Flink)
- Efficient data flow runtime on top of Hadoop/HDFS/YARN
 - ▶ Similar scalability and fail safety
 - ▶ More powerful data flow operators and optimization component
- Seamlessly integrates into existing Hadoop infrastructure:
 - ▶ can run side-by-side with Hadoop's TaskTrackers and DataNodes
 - ▶ can read data from Hadoop sources



Core Features

■ Enhanced Execution Engine

- ▶ Multiple data transformation:
 - Join, Cross, Union, Iterate, ...
- ▶ In-memory pipelining between operators
- ▶ Support for Iterative Algorithms
- ▶ Lazy Evaluation

■ Built-In Optimizer

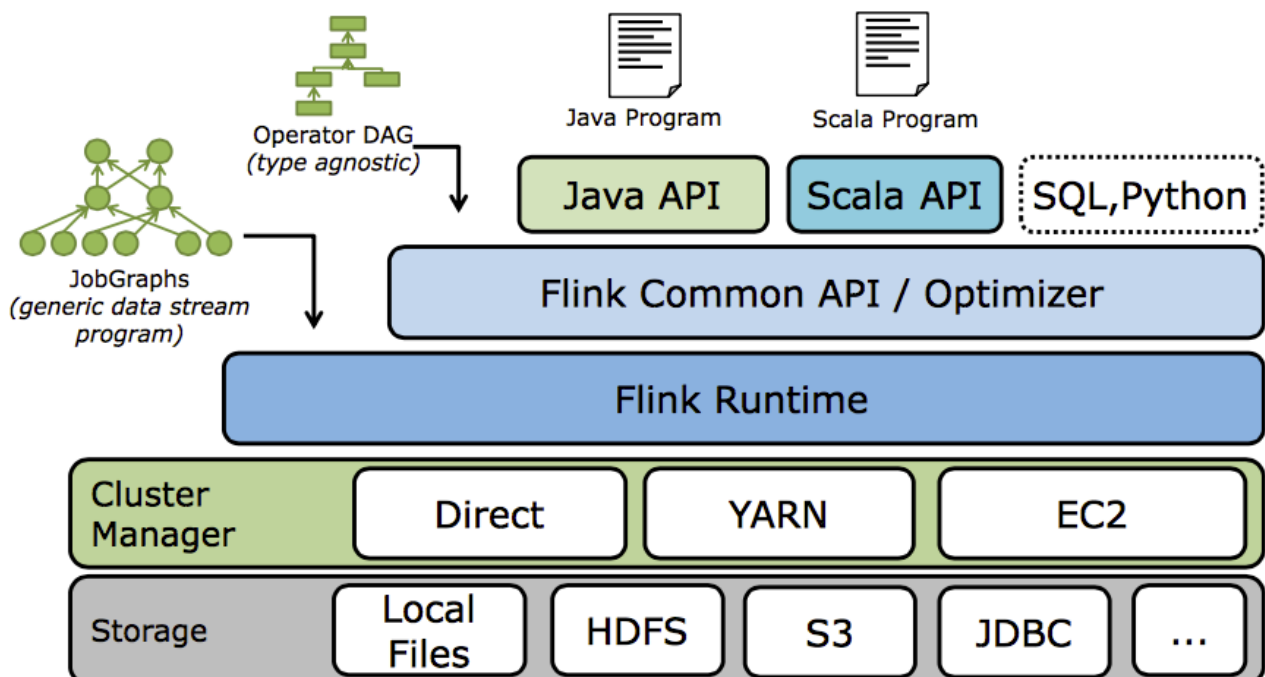
- ▶ Cost-based optimizer: execution strategy depending on inputs & ops
- ▶ For example: "Join" operator
 - Sort-merge-join, hash-join, broadcasting...
- ▶ Input Sampling to determine cardinalities

■ Enhanced APIs

- ▶ Support for Java and Scala



Flink System Stack



Source: http://ci.apache.org/projects/flink/flink-docs-release-0.8.1/internal_general_arch.html



Example: Scala API

■ Word count in Flink using Scala:

```
val input = TextFile(textInput)

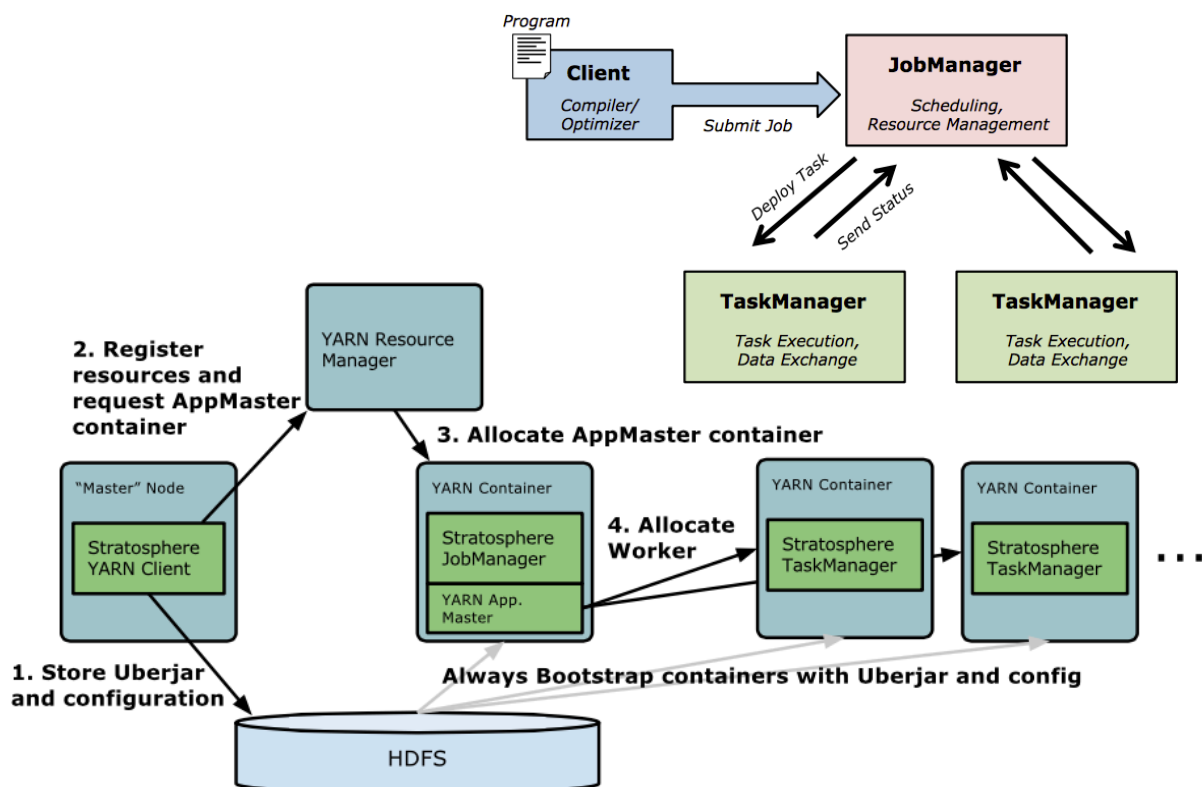
val words = input.flatMap { line => line.split(" ") }
val counts = words
  .map { (_, 1) }
  .groupByKey()
  .sum(1)

counts.print();

env.execute("Scala WordCount Example")
```



Flink on Yarn




Flink Programming Model

- Complex data analysis tasks specified as a data flow graph between parallelizable operators
 - ▶ Inherits aspects from MapReduce, but goes beyond it
- Aspects:
 - ▶ Data Model
 - ▶ Data Transformations
 - ▶ Sequential Data Flows
 - ▶ Iterative Data Flows



Flink: Data Model

- Inherent data model of Flink: **Tuple** (or record)
 - ▶ A tuple consists of an arbitrary number of (nested) data fields
 - ▶ Fields can be of a primitive data type, or a nested tuple structure



- General (primitive) Data Types (predefined)
 - ▶ **Integer values:** Byte, Short, Integer, Long
 - ▶ **Floating point values:** Float, Double
 - ▶ **Text values:** Character, String
 - ▶ **Special values:** Boolean, Null



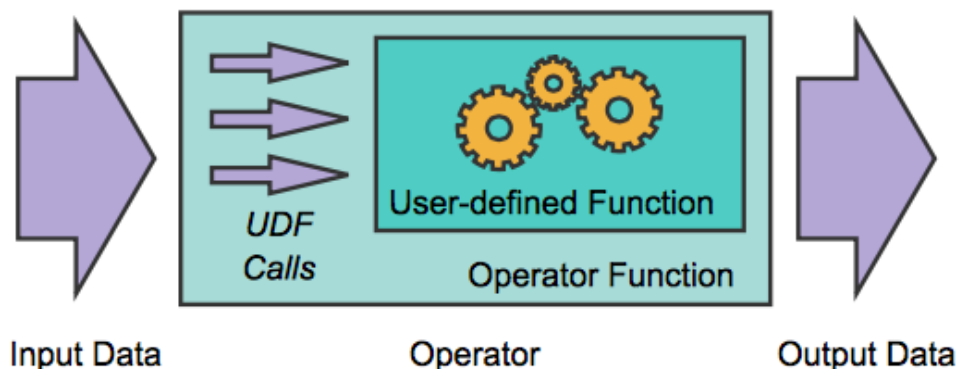
Flink Data Model (cont'd)

- Custom (User-defined) Data Types supported too
- **Value** types describe their serialization and deserialization manually
 - ▶ custom code for those operations can be implemented by means of the *org.apache.flinktypes.Value* interface with the methods *read* and *write*
 - ▶ Flink comes with pre-defined Value types that correspond to basic data types, and that in contrast to the basic data types are mutable. (ByteValue, ShortValue, IntValue, LongValue, FloatValue, DoubleValue, StringValue, CharValue, BooleanValue)
- A collection of tuples/records is called a **DataSet**
 - ▶ Operators receive one or more data sets as input and produce one new data set



Flink: Data Transformations

- programming model based on parallelizable operators
- An operator consists for two components
 - ▶ *user-defined function (UDF)* and
 - ▶ a *parallel operator function*.
 - The operator function parallelizes the execution of the user-defined function and applies the UDF on its input data.



Data Transformation (cont'd)

- Flink's programming model provides several parallelizable data transformations:

- ▶ *Map*,
- ▶ *Reduce* (including an optional *Combine*),
- ▶ *Filter*,
- ▶ *Aggregate*,
- ▶ *Join*,
- ▶ *Cross*,
- ▶ *CoGroup*,
- ▶ *Union*
- ▶ ...


[http://ci.apache.org/projects/flink/flink-docs-release-0.8/dataset_transformations.html]

- Map, Filter and Reduce operate on a single input
- Join, Cross, and CoGroup combine the data of two input
- Map, Filter, Join, and Cross operate on individual records, Reduce and CoGroup on groups of records



Map / FlatMap Operators

- *Record-at-a-Time* operator with one input
 - ▶ calls its user-defined function for each individual input record
- Same semantics as Hadoop MapReduce's Map:
 - ▶ The user function accepts a single record as input and can emit any number of records (0 to n for FlatMap, 1 for Map).
- Typical applications: filters or transformations




Reduce Operator

■ *Group-at-a-Time* operator with one input

- ▶ groups the tuples of its input on a so-called *Record Key* and hands each group into the user function
 - User function accepts list of tuples as input and can emit any number of tuples.
- ▶ Record Key: one or more fields of the input records that implement the *Key* interface.

■ Common application: aggregations



Reduce Optimisation: Combine

■ Reduce operator supports optional user-defined *Combine* function

- ▶ Optimisation for the case that the result of the Reduce operator can be computed from partial results that were computed on subgroups
- ▶ E.g. Sum:
The sum of a list of numbers, say $(1 + 2 + 3 + 4)$, can be computed by adding the sums of subgroups, e.g., $(1 + 2)$ and $(3 + 4)$.

■ Beneficial because it reduces the amount of data

- ▶ `Combine()` is called before data is transferred over the network to establish the full group for the final Reduce call

■ Note: Flink does not guarantee that the Combine function is actually executed.

- ▶ This is a cost-based decision done by Flink's optimizer.




Join Operator

■ *Record-at-a-Time* operator with two inputs

- ▶ requires the specification of record keys on both inputs
- ▶ operator function (equi-)joins both inputs on their record keys and hands matching pairs of records into the user function
- ▶ The user function accepts one record of each input and can emit any number of records

■ Typical application: equi join




Cross Operator

■ *Record-at-a-Time* operator with two inputs.

- ▶ In contrast to Join, Cross does not require the specification of a record key on any input.
- ▶ The operator function builds the Cartesian product of the records of both inputs and calls the user function for each pair of records.
- ▶ The user function accepts one record of each input and can emit any number of records.

■ Caution – very expensive operator that makes most sense on small inputs.




CoGroup Operator

■ *Group-at-a-Time* operator with two inputs.

- ▶ requires the specification of a record key for each input.
- ▶ Records of both inputs are grouped on their key. Groups with matching keys are handed together to the user function.
 - groups that have no matching group in the other input are given as a list to the user function and the other input list remains empty.
- ▶ The user function accepts one list of records for each input and can emit any number of records.

■ Application: Outer Join



Flink Optimizer

■ Flink has a **dedicated** optimizer

■ Goal: efficient execution plans for data processing tasks

- ▶ Estimates output sizes of each operator
- ▶ Based on cardinalities, decides on parallelism and local (in-memory) vs. shipping (via-network) pipelining strategy

■ Big challenge: UDFs

- ▶ Optimizer it is not aware of the internal semantics of an UDF
- ▶ Programming model allows to provide **annotations** for UDFs that explicitly provide the required information:
 - **Constant fields**: Lists the fields of a record that are NOT modified by the UDF (neither value or position in the record).
 - **All fields constant except**: Lists all fields that are modified by the UDF. All other fields are considered constant.
- ▶ Still, cardinalities and selectivities can only be guessed...



Data Flow

- The Flink programming model is based on data flow of parallelizable operators. The data flow is defined in the form of an *operator DAG* that consists of:
 - ▶ **Data Sources**
 - ▶ **Data Transformations**
 - ▶ **Data Sinks**



Data Sources

- A *data source* is the entry point of data into a data flow
 - ▶ Can connect to a variety of data stores, e.g. HDFS or relational or NoSQL database systems
 - ▶ Produces exactly one **DataSet**
- DS provide a generic interface called *InputFormat*
 - ▶ In general, InputFormats are also user-defined functions.
- Predefined InputFormats for a variety of sources including:
 - ▶ CSV files,
 - ▶ Row-delimited text files,
 - ▶ Binary files with constant record length,
 - ▶ Collection-based (creates Source by iterating over a collection)
 - ▶ Generic input DS eg. on external database sources (such as JDBC)




Data Sinks

- Data sinks are the exit point where data leaves a data flow.
- Similar to data sources, data sinks provide a generic interface called *OutputFormat* to write their input data set to a variety of data stores or streams.
 - ▶ An *OutputFormat* serializes records into a format, such as a textual or binary representation, and writes it to an interface outside of the system as for example a file system or database. *OutputFormats* are as well user-defined functions such that data can be written to any external data store.
- Flink provides *OutputFormats* for:
 - ▶ CSV files,
 - ▶ Row-delimited text files, and
 - ▶ Binary files.
 - ▶ Print to console



Data Flow Composition

- A data flow is composed of any number of data sources, transformations, and data sinks
 - ▶ cyclic data flows are disallowed
- Iterative data flows possible why **Iterate** operator



Example: Java API

- Cf. http://ci.apache.org/projects/flink/flink-docs-release-0.8/programming_guide.html
- <http://ci.apache.org/projects/flink/flink-docs-release-0.8/examples.html>
- <https://github.com/apache/flink/tree/master/flink-examples/flink-java-examples/src/main/java/org/apache/flink/examples/java>

```
public class WordCount {  
  
    public static void main (String[] args) {  
        // set up the execution environment  
        final ExecutionEnvironment env = ExecutionEnvironment.getExecutionEnvironment();  
  
        DataSet<..> source = ...  
  
        // Operations on the data set go here  
        DataSet<..> plan = source. ...  
  
        // define where to emit result of the execution plan  
        plan.writeAsCSV(outPath, "\n", " "); // alternative on-screen output: plan.print();  
  
        // initiate the actual execution  
        env.execute("WordCount Example");  
    }  
}
```



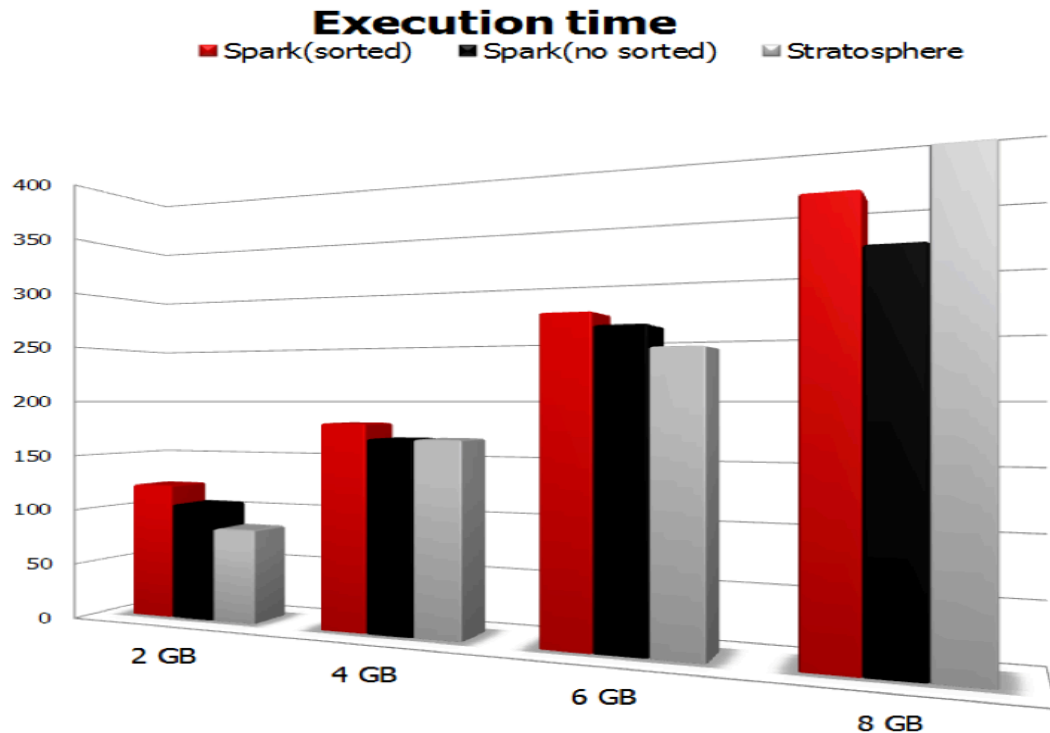
Example: Word Count

```
public class WordCount {  
  
    public static void main(String[] args) throws Exception {  
  
        // set up the execution environment  
        final ExecutionEnvironment env = ExecutionEnvironment.getExecutionEnvironment();  
  
        // get input data  
        DataSet<String> text = env.fromElements(  
            "To be, or not to be,--that is the question:--",  
            "Whether 'tis nobler in the mind to suffer",  
            "The slings and arrows of outrageous fortune",  
            "Or to take arms against a sea of troubles",  
        );  
  
        DataSet<Tuple2<String, Integer>> counts =  
            // split up the lines in pairs (2-tuples) containing: (word,1)  
            text.flatMap(new LineSplitter())  
            // group by the tuple field "0" and sum up tuple field "1"  
            .groupBy(0)  
            .aggregate(Aggregations.SUM, 1);  
  
        // emit result  
        counts.print();  
  
        // execute program  
        env.execute("WordCount Example");  
    }  
}
```

```
public class LineSplitter implements FlatMapFunction<String, Tuple2<String, Integer>> {  
  
    @Override  
    public void flatMap(String value, Collector<Tuple2<String, Integer>> out) {  
        // normalize and split the line into words  
        String[] tokens = value.toLowerCase().split("\\W+");  
  
        // emit the pairs  
        for (String token : tokens) {  
            if (token.length() > 0) {  
                out.collect(new Tuple2<String, Integer>(token, 1));  
            }  
        }  
    }  
}
```



Performance Comparison (WordCount)

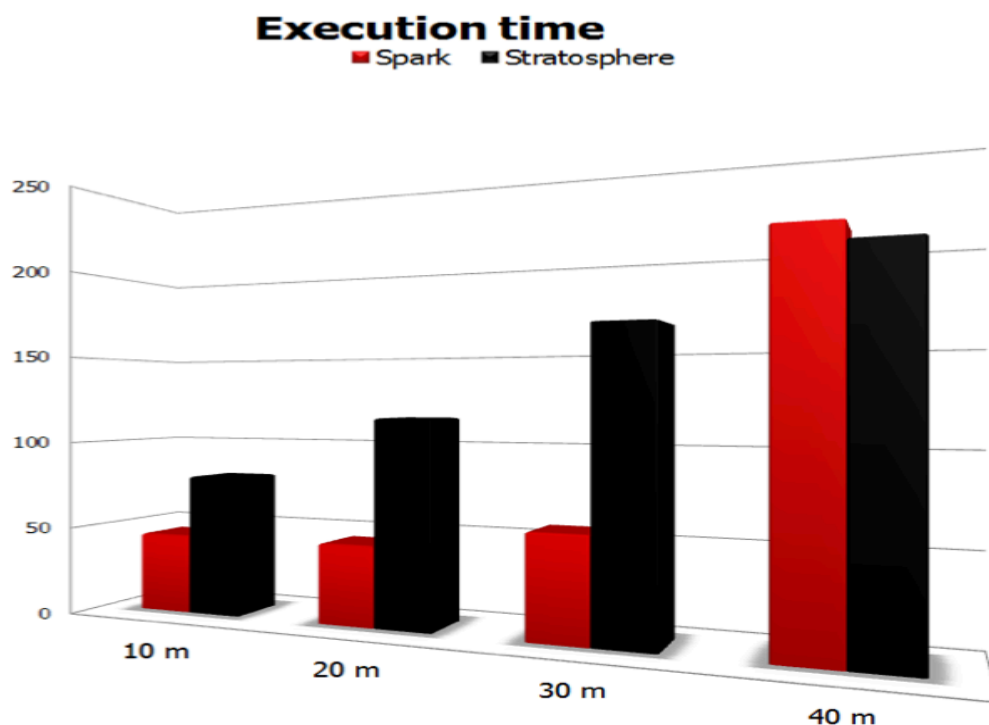


Word Count task on various data sizes in HDFS.

[Ze Ni, 2013]



Performance Comparison (K-Means)



K-means algorithm with 2 iterations on various data sizes (millions of points) in HDFS.

[Ze Ni, 2013]



Conclusions

- Several approaches exist for more sophisticated data analytics on top of Hadoop/HDFS infrastructure
 - ▶ strive to make data processing easier to program (more high-level)
 - ▶ support for iterative and streaming data processing
 - ▶ while keeping scalability and fault tolerance of MR
- Apache Spark:
 - ▶ Core idea: make distributed datasets a first-class primitive to provide a simple, efficient programming model for stateful data analytics
- Apache Flink:
 - ▶ Powerful data flow language on top of MapReduce
 - ▶ Integrates many ideas from traditional data processing, more focus on automated plan optimisation



References

- Spark / BDAS:
 - ▶ Matei Zaharia, Mosharaf Chowdhury, Michael J. Franklin, Scott Shenker, and Ion Stoica: "Spark: cluster computing with working sets". In *Proceedings of the 2nd USENIX conference on Hot topics in cloud computing*, HotCloud 2010.
 - ▶ M. Zaharia, M. Chowdhury, T. Das, A. Dave, J. Ma, M. McCauley, M. J. Franklin, S. Shenker, and I. Stoica: "Resilient distributed datasets: a fault-tolerant abstraction for in-memory cluster computing". In *USENIX NSDI*, 2012.
 - ▶ Ion Stoica: "Berkeley Data Analytics Stack (BDAS) Overview", Strata Conf., Feb 2013.
 - ▶ Matei Zaharia, Mosharaf Chowdhury, Justin Ma, Michael Franklin, Scott Shenker, Ion Stoica: "Spark: In-Memory Cluster Computing for Iterative and Interactive Applications", amplab.
 - ▶ Mike Franklin: "The Berkeley Data Analytics Stack: Present and Future", NICTA, Feb 2014.
- Flink / Stratosphere:
 - ▶ <http://flink.apache.org>
 - ▶ <http://stratosphere.eu/docs/>
 - ▶ Alexander Alexandrov, Dominic Battré, Stephan Ewen, Max Heimel, Fabian Hueske, Odej Kao, Volker Markl, Erik Nijkamp, and Daniel Warneke: "Massively parallel data analysis with PACTs on Nephele". *PVLDB* 3:1625–1628, Sept 2010.
- Ze Ni: "Comparative Evaluation of Spark and Stratosphere", master thesis, KTH Stockholm, 2013.



Lecture Outlook

■ Today:

- ▶ Working of Assignment 1 in the lab rooms in SIT with tutor feedback
- ▶ Submission tomorrow (Fri) by 6pm in Blackboard
- ▶ Self-Reflection Survey to be filled in too

■ Next Week: Details on Spark and FLink

- ▶ Including lab
- ▶ A2 to be published mid of Week 8/ early Week 9 too

■ Friday this week (28nd April): Data Centre Excursion

- ▶ Visit to one of the Equinix data centres in Alexandria (200 Bourke Rd)
- ▶ 2 tours @ 1-1:30pm and @ 1:45pm-2:30pm
=> **ENROL IN TOUR GROUPS IN Blackboard**



Usability Study of Cloud Computing Frameworks for Big Data Analytics

This semester, we also conduct a study of the usability and 'learnability' of cloud computing frameworks (such as Hadoop, Spark, Flink).

You can opt-in to use your assignment submissions and the self-reflection survey data for our study.

The study outcome will **not** be used for marking and everyone will need to do the same tasks; but if you opt-in, the **anonymised meta-data** of your assignment submissions and the answers from the self-reflective survey will be used for the study.

More details will be send by email later today.

