

COMP5349 –Cloud Computing

Week 8: Programming with Apache Spark and Flink

A/Prof. Uwe Roehm
School of Information Technologies



Self-Reflection Survey

- Please re-submit in eLearning as we had to re-create the self reflection survey on the weekend
- Please answer **all** questions in a single session
 - ▶ please review the instructions
 - ▶ if unsure, choose the middle option
- In the lab today:
 - ▶ Demos on Azure
 - ▶ Exercises with both Spark and Flink

Review Question 1

Which system uses Resilient Distributed Datasets (RDD)?

- **A:** Apache Flink
- **B:** Apache Hadoop
- **C:** Apache Spark



<http://tinyurl.com/mpkjmww>



Review Question 2

What is the main strength of Spark?

- **A:** in-memory processing
- **B:** fail safety
- **C:** pipelined processing



<http://tinyurl.com/knqekz6>



Review Question 3

What is the main strength of Flink?

- A: in-memory processing
- B: fail safety
- C: pipelined processing



<http://tinyurl.com/meqatdd>



Outline

- **Overview**
 - ▶ Spark & Flink Skeleton Programs
 - ▶ Wordcount Example
- **Spark and Flink Details**
 - ▶ RDD and Data Set Operations
- **Lambda Expression**
- **MR Design Patterns in Spark & Flink**
- **A Complete Example**



[based on slides by Dr Ying Zhou]



Spark & Flink Program Skeletons

■ Regular programs in Java / Scala / Python

1. Initialise the runtime **environment**
2. Load or create **source** data
3. Specify the **data transformations**
 - ▶ This might include usage of self-defined UDFs
4. Specify where the **output** should go
5. **Execute**



Comparison

- In the following, each step with either Spark or Flink
- First Java, but also all examples in Python available



Step 1: Initialisation

■ Spark:

- Every Spark application represented as a driver program
 - ▶ In Java, the driver program is just a normal java class with a main method
- Every Spark application needs to create a **SparkContext** object, which is used to access a cluster. In Java language The **SparkContext** object is of type **JavaSparkContext**

■ Flink:

- In Java, a Flink main program is a normal Java class with a main() method
- Every Flink application needs to create an ExecutionEnvironment object, which is used to access the cluster. In Java, this object is of type **ExecutionEnvironment**



Spark – Driver Initialisation

- A Spark program is just a regular **main** program that creates **SparkContext** object, which is used to access a cluster.
 - ▶ In Java, this object is of type **JavaSparkContext**
 - ▶ In Python, this object is of type ...
- The spark context provides methods for
 - ▶ Data input
 - ▶ Data output
 - ▶ ...
- The data processing steps are defined using **RDDs**



WordCount in Spark (Java)

```
import java.util.Arrays;
import org.apache.spark.SparkConf;
import org.apache.spark.JavaSparkContext;
import org.apache.spark.JavaRDD;
import org.apache.spark.JavaPairRDD;

public class WordCountExample {
    public static void main(String[] args) {
        SparkConf conf = new SparkConf().setAppName("WordcountExample")
                                         .setMaster("...");
        JavaSparkContext sc = new JavaSparkContext(conf);

        JavaRDD<String> textFile = sc.textFile("hdfs://...");
        JavaPairRDD<String, Integer> counts =
            textFile.flatMap(s -> Arrays.asList(s.split(" ")).iterator())
                    .mapToPair(word -> new Tuple2<>(word, 1))
                    .reduceByKey((a, b) -> a + b);

        counts.saveAsTextFile("hdfs://...");
    }
}
```

[Cf.: <http://spark.apache.org/examples.html>]



WordCount in Spark (Python)

```
from pyspark import SparkConf, SparkContext

myconf = (SparkConf().setMaster("...")
          .setAppName("WordcountExample")
          .set("spark.executor.memory", "1g")
          .pyFiles(['wordcount.py']))
sc = SparkContext(myconf)

text_file = sc.textFile("hdfs://...")
counts = text_file.flatMap(lambda line: line.lower.split(" ")) \
               .map(lambda word: (word, 1)) \
               .reduceByKey(lambda a, b: a + b)
counts.saveAsTextFile("hdfs://...")
```

[Cf.: <http://spark.apache.org/examples.html>]



Spark– Linking and Execution

- Need to add additional libraries to the program
- Spark Java:
 - ▶ spark-assembly-1.6.1-hadoop2.6.0.jar
 - ▶ spark-core_2.10-0.9.0-incubating.jar



PySpark Shell

- With PySpark, there is also the possibility to use the Python shell to explore data interactively and as a simple way to learn the API:

```
./bin/pyspark
>>> words = sc.textFile("hdfs://...")
>>> words.filter(lambda w: w.startswith("spar")).take(5)
[u'spar', u'sparable', u'sparada', u'sparadrap', ...]
>>> help(pyspark) # Show all pyspark functions
```

- By default, the bin/pyspark shell creates a SparkContext 'sc' that runs applications locally on a single core. To connect to a non-local cluster, or use multiple cores, set the MASTER environment variable.

[<https://spark.apache.org/docs/0.9.1/python-programming-guide.html>]



Flink – Program Initialisation

- A Flink program is just a regular **main** program that creates an Flink “execution environment” (an `ExecutionEnvironment` object), which is used to access the cluster
 - ▶ In Java, this object is of type `ExecutionEnvironment`
 - ▶ In Python, this object is of type `flink.plan.Environment`
- The execution environment provides methods for
 - ▶ Data input
 - ▶ Data output
 - ▶ `execute()`: actually executes the defined Flink workflow in the cluster
- The processing workflow is defined using `DataSet` objects and its API to manipulate data sets



WordCount in Flink (Java)

```
public class WordCountExample {
    public static void main(String[] args) {
        final ExecutionEnvironment env = ExecutionEnvironment.getExecutionEnvironment();
        DataSet<String> data = env.readTextFile("hdfs://...");

        DataSet<Tuple2<String,Integer>> wordCounts =
            data.flatMap(new Tokenizer())
                .groupBy(0) // group by tuple field "0"
                .sum(1);    // sum up tuple field "1"
        wordCounts.writeAsCsv("hdfs://...", "\n", "");
        env.execute("Word Count Example");
    }

    public static class Tokenizer
        implements FlatMapFunction<String, Tuple2<String,Integer>> {
        @Override public void flatMap(String line,
                                      Collector<Tuple2<String, Integer>> out) {
            for (String word : line.toLowerCase().split(" "))
                out.collect(new Tuple2<String, Integer>(word, 1));
        }
    }
}
```

[Cf.: <https://ci.apache.org/projects/flink/flink-docs-release-1.2/dev/batch/examples.html>]



WordCount in Flink (Python)

```
from flink.plan.Environment import get_environment
from flink.functions.GroupReduceFunction import GroupReduceFunction

class Sum (GroupReduceFunction):
    def reduce(self, iterator, collector):
        word, count = iterator.next();
        count += sum([x[1] for x in iterator])
        collector.collect((word,count))

env = get_environment()
data= env.read_text("hdfs://...");

data.flat_map(lambda x,c: [(word,1) for word in x.lower().split()]) \
    .group_by(0) \
    .reduce_group(Sum(), combinable=True) \
    .write_csv("hdfs://...")

env.execute(local=True)
```

[Cf.: <https://ci.apache.org/projects/flink/flink-docs-release-1.2/dev/batch/python.html>]



Flink – Linking and Execution

- need to include the corresponding Flink library in our project
- simplest way is via a Maven (build tool) project template
- Flink Java:

```
mvn archetype:generate /
    -DarchetypeGroupId=org.apache.flink/
    -DarchetypeArtifactId=flink-quickstart-java /
    -DarchetypeVersion=0.8.1
```

- Flink Python:

- ▶ Apart from setting up Flink, no additional work is required. The flink package, along with the plan and optional packages are automatically distributed among the cluster via HDFS when running a job. Python 2.7 or 3.4 should be used.



Lessons Learned

- Flink's Python binding is not as mature as its Java API
 - ▶ Eg. more manual function definitions needed
- Some operations are pre-defined, but in some we need to introduce and use own functions
 - ▶ Cf. in Python example the reduce function before where we need to define an appropriate **sum** aggregator first
 - ▶ Or in the Java code, the **tokenizer** is user-defined



Outline

- Overview
 - ▶ Spark & Flink Skeleton Programs
 - ▶ Wordcount Example
- Spark and Flink Details
 - ▶ RDD and Data Set Operations
- Lambda Expression
- MR Design Patterns in Spark & Flink
- A Complete Example



More on Spark

- The following is another example that relates to the MovieLens data set used in the tutorial

```
public class MovieLens {  
  
    public static void main(String[] args) {  
        SparkConf conf = new SparkConf();  
        // we can set lots of information here  
        conf.setAppName("Movie Lens Application");  
        JavaSparkContext sc = new JavaSparkContext(conf);  
  
        // create some RDDs  
        JavaRDD<String> ratingData = sc.textFile(inputDataPath+"u.data"),  
            movieData = sc.textFile(inputDataPath + "u.item");  
        // lots and lots of RDD operations  
        JavaPairRDD<String, Integer> ratingExtraction =  
            ratingData.mapToPair(s -> {  
                String[] values = s.split("\t");  
                return new Tuple2<String, Integer>(values[1], Integer.parseInt(values[2]));  
            });  
        ...  
    }  
}
```



Spark: Resilient Distributed Datasets

- RDDs are created by
 - ▶ Parallelizing an existing collection
 - ▶ Referencing a dataset in an external storage system, such as HDFS
- RDDs have partitions
 - ▶ Based on source file partition (such as blocks of HDFS files)
 - ▶ Or created during transformation, repartition

```
//parallelizing existing collection  
List<Integer> data = Arrays.asList(1, 2, 3, 4, 5);  
JavaRDD<Integer> distData = sc.parallelize(data);  
  
//referencing a data set in HDFS  
JavaRDD<String> ratingData =  
    sc.textFile("hdfs://ip-10-171-118-84.ec2.internal:8020/share/ml/u.data");
```



Handling Key Value Pairs in Spark

- Spark have special RDD types to handle data in key value format
 - ▶ In Java, key-value pairs are represented using the `scala.Tuple2` class from the `Scala` standard library.
 - ▶ The RDDs of key-value pair are represented by the `JavaPairRDD` type
 - `JavaPairRDD<String, Integer>`
- `JavaPairRDD`s are constructed from norm `JavaRDD`s



Spark RDD Operations

- Transformation
 - ▶ create a new dataset from an existing one
 - ▶ Eg. `map(func)`, `flatMap(func)`, `mapToPairs(func)`, `reduceByKey(func)`
- Action
 - ▶ return a value to the driver program after running a computation on the data set
 - ▶ Eg. `count()`, `first()`, `collect()`, `saveAsTextFile(path)`
- Most RDD operations take one or more functions as parameter
 - ▶ Most of them can be viewed as higher order functions
- Spark has strong functional programming flavour!



Flink DataSet Operations

- Functionality similar to the RDDs in Spark
 - ▶ Cf. documentation
 - ▶ No in-memory caching though, but inherently pipelined processing
- Most DataSet operations take one or more functions as parameter
 - ▶ Most of them can be viewed as higher order functions
- Flink also has a strong functional programming flavour!



Outline

- **Overview**
 - ▶ Spark & Flink Skeleton Programs
 - ▶ Wordcount Example
- **Spark and Flink Details**
 - ▶ RDD and Data Set Operations
- **Lambda Expression**
- **MR Design Patterns in Spark & Flink**
- **A Complete Example**



Anonymous Functions

- Functional argument can be compactly expressed using anonymous functions
 - ▶ E.g: `map (a-> 2*a)`
- **Lambda expressions** are a way to express anonymous functions in a programming language
- Nearly all functional languages and most script languages support anonymous function
- Java supports anonymous functions beginning with Java 8 in the form of lambda expression
 - ▶ **Function** interface is used in previous versions



Lambda Expressions in Java

```
// create an RDD from local file system as a collection of strings
JavaRDD<String> lines = sc.textFile("data.txt");
// create an RDD of Integer using the map transformation
// the transformation take a line as string and returns the number
// of characters
JavaRDD<Integer> lineLengths = lines.map(s -> s.length());
// call the reduce action to find total number of characters in the file
int totalLength = lineLengths.reduce((a, b) -> a + b);
```

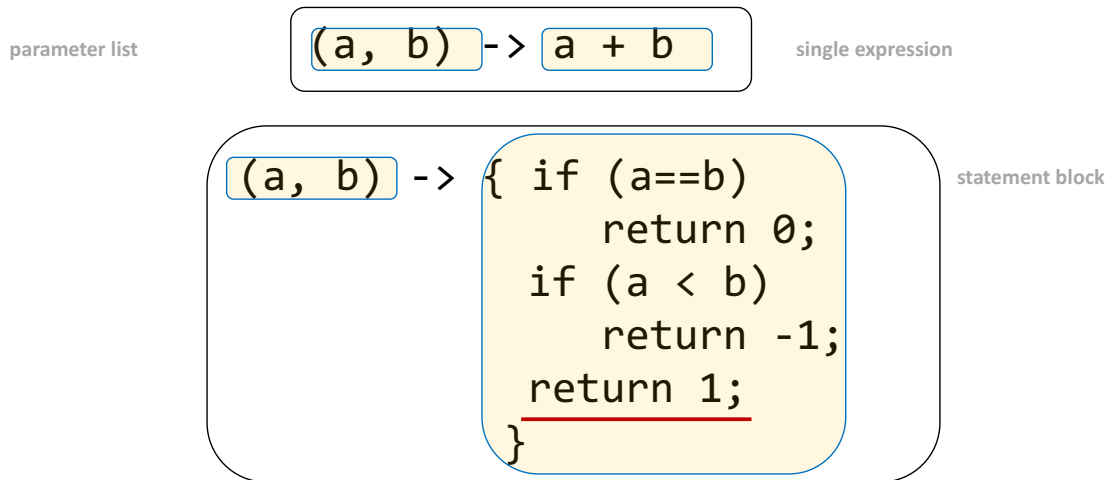
All transformations in Spark or Flink are *lazy*, in that they do not compute their results right away. Instead, they just remember the transformations applied to some base dataset (e.g. a file). The transformations are only computed when an action requires a result to be returned to the driver program.

Adding a line `lineLengths.persist();` before reduce action would cause `lineLengths` to be saved in memory for later use



Lambda Expression in Java

- A Lambda Expression consists of the following
 - ▶ A comma-separated lists of formal parameters enclosed in parentheses
 - ▶ The arrow token , ->
 - ▶ A body, which consists of a single expression or a statement block



RDD Operation Functional Argument

String -> int

```
JavaRDD<Integer> lineLengths = lines.map(s -> s.length());

int totalLength = lineLengths.reduce((a, b) -> a + b);
```

(int, int) -> int

map

<R> JavaRDD<R> map(Function<T,R> f)
Return a new RDD by applying a function to all elements of

public interface Function<T1,R>
extends java.io.Serializable

Method Summary

Methods

Modifier and Type	Method and Description
R	call(T1 v1)

reduce

T reduce(Function2<T,T,T> f)
Reduces the elements of this RDD using the specified commutative and associative binary operator.

public interface Function2<T1,T2,R>
extends java.io.Serializable

A two-argument function that takes arguments of type T1 and T2 and returns an R.

Method Summary

Methods

Modifier and Type	Method and Description
R	call(T1 v1, T2 v2)



Anonymous Functions with Flink

- The previous examples were using Java and the Spark API
- Same concept (lambda expressions) are also supported natively in Python
- Flink API also relies heavily on lambda expressions
 - ▶ Please see documentation and lab for details



Outline

- **Overview**
 - ▶ Spark & Flink Skeleton Programs
 - ▶ Wordcount Example
- **Spark and Flink Details**
 - ▶ RDD and Data Set Operations
- **Lambda Expression**
- **MR Design Patterns in Spark & Flink**
- **A Complete Example**



MR Design Patterns with Spark or Flink

- In Week 5, we covered a series of MR design patterns
- In the following, want to briefly run through how those can be represented in Spark or Flink
- Input Transformation
 - ▶ In MR: Part of the map() function
 - ▶ In Spark: New RDD based on initial RDD;
 - ▶ In Flink: New DataSet based on initial Data Set



MR Design Pattern: Filtering

- Filtering input data
 - ▶ In MR: Part of the map() function
 - ▶ In Spark: `RDD.filter(predicate)` operation
 - Only entries where the *predicate* function returns TRUE are kept
 - ▶ In Flink: ...
 - Only entries where the *predicate* function returns TRUE are kept
- Example:
 - ▶ Spark:

```
JavaRDD<Movie> movies = sc. ...  
JavaRDD<Movie> result = movies.filter( m -> 1 <= m.id <= 100);
```
 - ▶ Flink:

```
DataSet<String> data = env.readTextFile("hdfs://...");  
DataSet<String> result = data.filter( line -> ! line.contains("error") );
```



MR Design Pattern: Aggregation

■ Aggregating values

- ▶ In MR: implemented with the `reduce()` function
- ▶ In Spark: `reduceByKey()`
- ▶ In Flink: sequence of `groupByKey()` followed by `reduce()`
- ▶ both systems offer predefined aggregation functions such as `sum()` or `count()`

■ Example:

▶ Spark:

```
JavaPairRDD<String, Integer> counts =  
    textFile.flatMap( ... )  
              .mapToPair(word -> new Tuple2<>(word, 1))  
              .reduceByKey((a, b) -> a + b);
```

▶ Flink:

```
DataSet<Integer> textfile.flatMap( { ... return (word,1) } )  
                          .groupBy( 0 )  
                          .sum(1);
```



MR Design Pattern: Join

■ Joining two input data sets

- ▶ In MR: pair of `map()` and `reduce()` function with appropriate keys
- ▶ In Spark: `join()` method of PairRDDs (we need a key per value)
 - performs a hash join across the cluster with resulting `<k,(v1, v2)>` pairs where `(k, v1)` was in the first, and `(k, v2)` in the second RDD
- ▶ In Flink: `join()` method on Data Sets
- ▶ both offer various sub-forms of joins, such as cross-join (Flink), or left or right outer join (Spark)

■ Example:

▶ Spark

```
JavaPairRDD<> input1, input2;  
JavaPairRDD<> result = input1.join(input2)
```

▶ Flink:

```
DataSet<> input1, input2;  
DataSet<> result = input1.join(input2)  
                        .where(0) // key of first input (tuple field 0)  
                        .equalTo(1); // key of 2nd input (tuple field 1)
```



Combiner, Partitioner, Composite Keys

■ Combiner

- ▶ Spark PairedRDDs offer a `combineByKey()` method

■ Partitioner

- ▶ Spark RDDs can get their `partitioner` class attribute overridden
- ▶ Spark PairedRDDs have a `partitionBy()` method
- ▶ Flink DataSets support `partitionByHash()`, `partitionByRange()` and `partitionCustom()` methods

■ Composite Keys

- ▶ In Spark:
- ▶ In Flink: quite natural via Tuple specification

[cf. <https://spark.apache.org/docs/0.6.2/api/core/spark/RDD.html>]

[cf. <https://ci.apache.org/projects/flink/flink-docs-release-1.2/dev/batch/index.html>]



Combining two data values

■ Problem: Both Flink and Spark use **Lazy Evaluation**

- ▶ You cannot use single result values of one computation inside another, because it has not been computed yet, but is just a placeholder for a 'sub-query'

■ Two approaches:

- ▶ 1. Cross-Join
- ▶ 2. Broadcast variables (Flink)



Outline

■ Overview

- ▶ Spark & Flink Skeleton Programs
- ▶ Wordcount Example

■ Spark and Flink Details

- ▶ RDD and Data Set Operations

■ Lambda Expression

■ MR Design Patterns in Spark & Flink

■ A Complete Example

- ▶ Various transformations



Sample Program

■ Two data sets stored as txt files

■ Movies (mid, title, genres)

- ▶ Sample data:
▶ 1, Toy Story (1995), Adventure|Animation|Children|Comedy|Fantasy

■ Ratings (uid, mid, rating, timestamp)

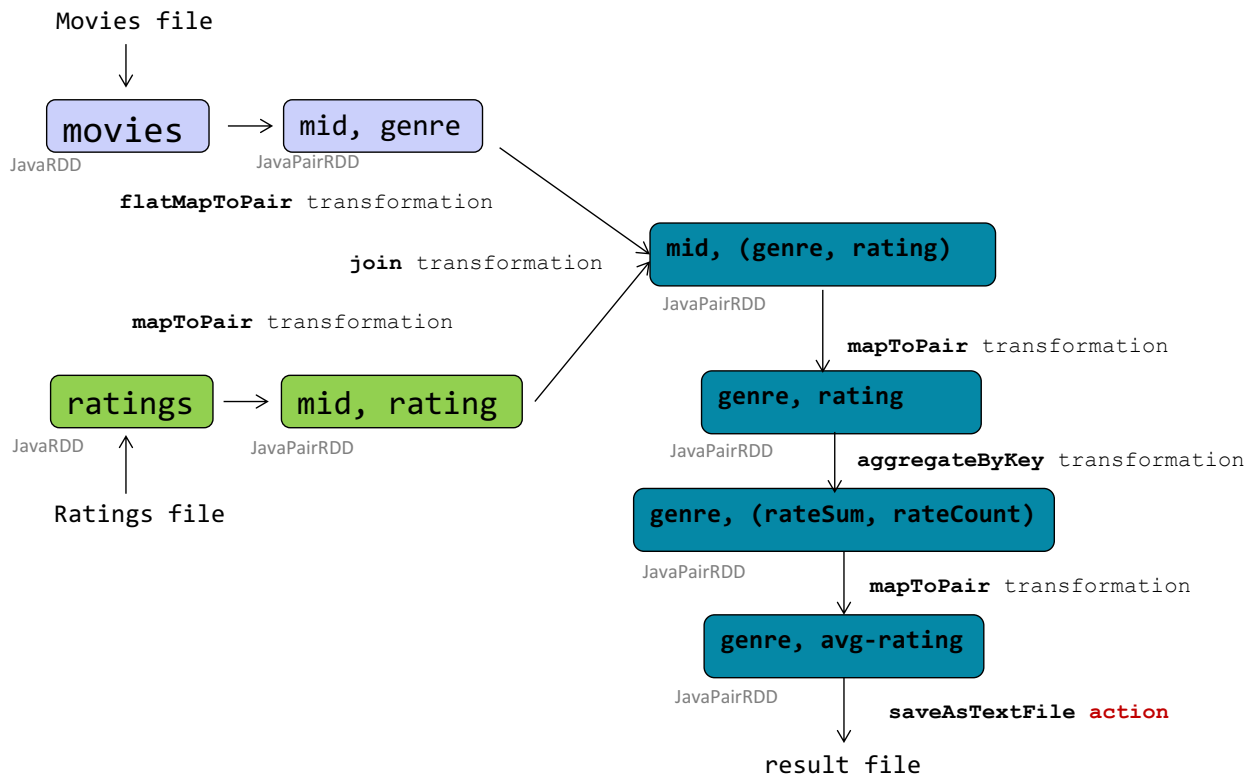
- ▶ Sample data:
▶ 1, 253, 3.0, 900660748

■ We want to find out the average rate for each genre

- ▶ We would join the two data sets on movie id (**mid**) and keep only the genre and rating data, we then group the rating data based on genre and find the average for each genre.



Spark: RDD Operation Design



Spark Application Code Snippet: I

```

public class MovieLens {

    public static void main(String[] args) {
        String inputDataPath = args[0], outputPath = args[1];
        //Prepare Spark Context
        SparkConf conf = new SparkConf();
        conf.setAppName("Movie Lens Application");
        JavaSparkContext sc = new JavaSparkContext(conf);

        //read ratings and movie data as JavaRDD<String>
        JavaRDD<String> ratingData = sc.textFile(inputDataPath+"ratings.csv"),
            movieData = sc.textFile(inputDataPath + "movies.csv");

        //convert ratingData to a key value pair RDD (mid, rating)
        //the text file format is: userId,movieId,rating,timestamp

        JavaPairRDD<String, Float> ratingExtraction = ratingData.mapToPair(s ->
        {
            String[] values = s.split(",");
            return
                new Tuple2<String, Float>(values[1],Float.parseFloat(values[2]));
        });
    }
}
  
```

Lambda expression

In Java, a key value pair is wrapped as a type **Tuple2**. The **mapToPair** transformation takes a function of type **PairFunction** as parameter, this function takes a parameter and returns a **Tuple2** object

mapToPair

```

<K2,V2> JavaPairRDD<K2,V2> mapToPair(PairFunction<T,K2,V2> f)
  
```

Return a new RDD by applying a function to all elements of this RDD.



Spark Application Code Snippet: II

```
//convert movieData to a key value pair RDD of format (mid, genre)
//the text file format is: movieID,title,genre1|genre2|genre3
```

```
JavaPairRDD<String,String> movieGenres = movieData.flatMapToPair(s->{
    String[] values = s.split(",");
    String movieID = values[0];
    int length = values.length;
    ArrayList<Tuple2<String,String>> results = new
        ArrayList<Tuple2<String,String>>();
    if (length >= 3) { // genre list is present
        String[] genres = values[length - 1].split("\\|");
        for (String genre: genres) {
            results.add(new Tuple2<String, String>(movieID, genre));
        }
    }
    return results;
});
```

Lambda expression

A **map** or **mapToPair** transformation always expects a one-to-one mapping between the input and output.

If you need to filter the input by some condition and only return the value that satisfy the condition, you should use **filter** transformation.

If an input would produce 0 or more output, you should use **flatMap** or **flatMapToPair** transformation. The **flatMapToPair** transformation takes a function of type **PairFlatMapFunction** as parameter. This function takes a parameter and returns an **Iterable** collection of **Tuple2** objects



mapToPair and flatMapToPair

mapToPair

```
<K2,V2> JavaPairRDD<K2,V2> mapToPair(PairFunction<T,K2,V2> f)
```

Return a new RDD by applying a function to all elements of this RDD.

```
public interface PairFunction<T,K,V>
    extends java.io.Serializable
```

A function that returns key-value pairs (Tuple2<K, V>), and can be used to construct PairRDDs.

Method Summary

Methods

Modifier and Type	Method and Description
scala.Tuple2<K,V>	call(T t)

flatMapToPair

```
<K2,V2> JavaPairRDD<K2,V2> flatMapToPair(PairFlatMapFunction<T,K2,V2> f)
```

Return a new RDD by first applying a function to all elements of this RDD, and then flattening the results.

```
public interface PairFlatMapFunction<T,K,V>
    extends java.io.Serializable
```

A function that returns zero or more key-value pair records from each input record. The key-value pairs are represented as scala.Tuple2 objects.

Method Summary

Methods

Modifier and Type	Method and Description
Iterable<scala.Tuple2<K,V>>	call(T t)



Spark Application Code Snippet: III

```
//join the two RDDs to find the ratings for each genre
//join function performs an inner join
//The result RDD would have the following format
//(movieID, (genre, rating))

JavaPairRDD<String, Tuple2<String,Float>> joinResults =
    movieGenres.join(ratingExtraction);

//Join is based on movieID, which is not useful in our calculation
//We only want to retain the value which is (genre, rating) and
//convert it to a PairRDD
JavaPairRDD<String, Float> joinResultsNoID = joinResults.values().mapToPair(v->v);
```

An RDD of `Tuple2<String,Float>`

A `Tuple2<String,Float>`

Return the same `Tuple2<String,Float>`



Spark Application Code Snippet: IV

```
JavaPairRDD genreRatingAvg = joinResultsNoID.aggregateByKey(
    new Tuple2<Float, Integer> (0.0f,0),
    1,
    (v1,v2)-> new Tuple2<Float, Integer> (v1._1+ v2, v1._2+1),
    (v3,v4) -> new Tuple2<Float,Integer> (v3._1 + v4._1, v3._2 + v4._2))
    .mapToPair(
        t -> new Tuple2(t._1, (t._2._1 * 1.0 / t._2._2))
    );

// this is an action
genreRatingAvg.saveAsTextFile(outputDataPath + "latest.rating.avg.per.genre");
sc.close();
}
```

There are a few transformations similar to the **reducer** in MapReduce framework

groupByKey groups the values for each key. This is like the step of preparing input for reducer in MapReduce framework

reduceByKey merge the values for each key using a given reduce function. This will also perform the merging locally on each "mapper" before sending results to a reducer, similarly to a "combiner" in MapReduce. But the type of the merged value should be the same as the type of the input value!

foldByKey is similar to **reduceByKey** except that you can supply a natural zero value.

aggregateByKey is more general than **reduceByKey**. It allows the merged value to have different type of the input value. It takes at least a natural zero value and two functions as parameter.

All above transformations can take extra parameter to indicate the number of partition, or a partitioner object. This is like specifying the number of **reducers** in MapReduce, or specifying a customized partitioner.




aggregateByKey Transformation

aggregateByKey

```
public <U> JavaPairRDD<K,U> aggregateByKey(U zeroValue,
    int numPartitions,
    Function2<U,V,U> seqFunc,
    Function2<U,U,U> combFunc)
```

Aggregate the values of each key, using given combine functions and a neutral "zero value". This function can return a different result type, U, than the type of the values in this RDD, V. Thus, we need one operation for merging a V into a U and one operation for merging two U's, as in `scala.TraversableOnce`. The former operation is used for merging values within a partition, and the latter is used for merging values between partitions. To avoid memory allocation, both of these functions are allowed to modify and return their first argument instead of creating a new U.



Spark Application Code Snippet: IV (revisit)

```
JavaPairRDD genreRatingAvg = joinResultsNoID.aggregateByKey(
    new Tuple2<Float, Integer> (0.0f,0),
    1,
    (r,v)-> new Tuple2<Float, Integer> (r._1+ v2, r._2+1),
    (v1,v2) -> new Tuple2<Float,Integer> (v1._1 + v2._1, v1._2 + v2._2))
    .mapToPair(
        t -> new Tuple2(t._1, (t._2._1 * 1.0 / t._2._2))
    );

// this is an action
genreRatingAvg.saveAsTextFile(outputDataPath + "latest.rating.avg.per.genre");
sc.close();
}
```

we want to use **aggregateByKey** to find out the rating sum and rating count per genre. the sum and count can be stored in a **Tuple2** object of float and integer, initialized to 0.0 and 0


For each rating value, we increment the sum and the count respectively

To merge two intermediate results, the partial sums are added up and the partial count added up


A **mapToPair** transformation is applied to calculate the average using the final sum and count



aggregateByKey Running Example



Partition 1



Partition 2

A hashmap `accumulators` is created to store result

```
Read ("Action", 5.0), "Action" is a new key
accumulators["Action"] = <0.0,0>
accumulators["Action"] = <0.0+5.0,0+1> = <5.0,1> //call seqFunc

Read ("Drama", 3.5), "Drama" is a new key
accumulators["Drama"] = <0.0,0>
accumulators["Drama"] = <0.0+3.5,0+1> = <3.5,1> //call seqFunc

Read ("Action", 4.0), "Action" is an old key
accumulators["Action"] = <5.0+4.0,1+1> = <9.0,2> //call seqFunc
```

("Action", <9.0,2>)
("Drama", <3.5,1>)

A hashmap `accumulators` is created to store result

```
Read ("Action", 3.0), "Action" is a new key
accumulators["Action"] = <0.0,0>
accumulators["Action"] = <0.0+3.0, 0+1> = <3.0,1>

Read ("Drama", 4.5), "Drama" is a new key
accumulators["Drama"] = <0.0,0>
accumulators["Drama"] = <0.0+4.5, 0+1> = <4.5,1>

Read ("Drama", 3.0), "Drama" is a new key
accumulators["Drama"] = <0.0,0>
accumulators["Drama"] = <4.5+3.0, 1+1> = <7.5,2>
```

("Action", <3.0,1>)
("Drama", <7.5,2>)

Merge by calling `combFunc`

("Action", <12.0,3>)
("Drama", <11.0,3>)

```
new Tuple2<Float, Integer> (0.0f,0) // init value
(v1,v2)-> new Tuple2<Float, Integer> (v1._1+ v2, v1._2+1) // seqFunc
(v3,v4) -> new Tuple2<Float,Integer> (v3._1 + v4._1, v3._2 + v4._2) // combFunc
```



References

- Zaharia, Matei, et al. "Resilient distributed datasets: A fault-tolerant abstraction for in-memory cluster computing." *Proceedings of the 9th USENIX conference on Networked Systems Design and Implementation*. USENIX Association, 2012.
 - ▶ https://www.cs.berkeley.edu/~matei/papers/2012/nsdi_spark.pdf
- Spark Document, Spark Programming Guide
 - ▶ <https://spark.apache.org/docs/1.2.0/programming-guide.html#parallelized-collections>
 - ▶ Submitting Applications in Spark: <https://spark.apache.org/docs/1.2.1/submitting-applications.html>
- Flink 1.2 Documentation, Flink Programming Guide
 - ▶ <https://ci.apache.org/projects/flink/flink-docs-release-1.2/dev/batch/index.html>
- Sandy Ryza, How-to: Tune Your Apache Spark Jobs (part 1) published on March 09, 2015
 - ▶ <http://blog.cloudera.com/blog/2015/03/how-to-tune-your-apache-spark-jobs-part-1/>
- Sandy Ryza, How-to: Tune Your Apache Spark Jobs (part 2) published on March 30, 2015
 - ▶ <http://blog.cloudera.com/blog/2015/03/how-to-tune-your-apache-spark-jobs-part-2/>
- AWS Webinar about Spark on AWS:
 - <https://aws.amazon.com/webinars/anz-webinar-series/>

