

COMP9120

Week 4: Relational Algebra and SQL

Semester 2, 2016

(Ramakrishnan/Gehrke – Chapter 4.2 & 5;

Kifer/Bernstein/Lewis – Chapter 5;

Ullman/Widom – Chapter 2.4 & 6)

Based on material by Dr. Bryn Jeffries





COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1969

WARNING

This material has been reproduced and communicated to you by or on behalf of the University of Sydney pursuant to Part VB of the Copyright Act 1968 (**the Act**).

The material in this communication may be subject to copyright under the Act. Any further copying or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice.

Flashback Question

<i>title</i>	<i>year</i>	<i>length</i>	<i>genre</i>
Gone With the Wind	1939	231	drama
Star Wars	1977	124	sciFi
Wayne's World	1992	95	comedy

Equivalent Relations?

<i>year</i>	<i>genre</i>	<i>title</i>	<i>length</i>
1977	sciFi	Star Wars	124
1992	comedy	Wayne's World	95
1939	drama	Gone With the Wind	231

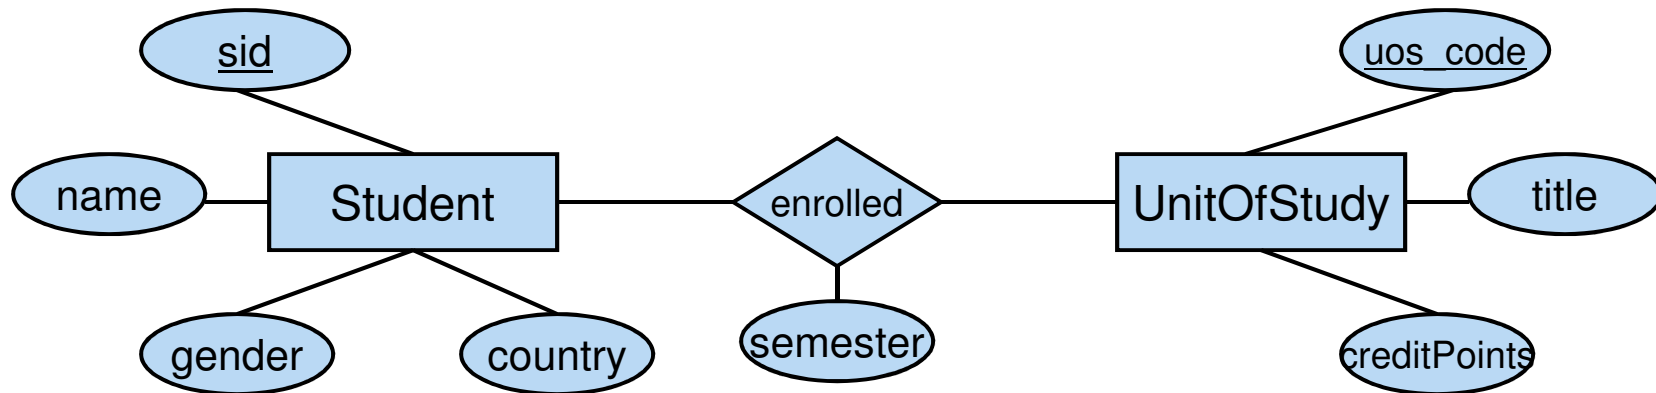
J. Ulman & J Widom, A First Course in Database Systems, 2008

› **Foundations of Declarative Querying**

- Relational Algebra
- Six basic operations
- Join operation
- Composition and equivalence rules

› **Introduction to SQL**

- Overview
- Basic SQL Queries
- Joins Queries
- Aggregate Functions and Set Operations



Student			
<u>sid</u>	name	gender	country
1001	Ian	M	AUS
1002	Ha Tschi	F	ROK
1003	Grant	M	AUS
1004	Simon	M	GBR
1005	Jesse	F	CHN
1006	Franziska	F	GER

Enrolled		
<u>sid</u>	<u>uos_code</u>	semester
1001	COMP5138	2005-S2
1002	COMP5702	2005-S2
1003	COMP5138	2005-S2
1006	COMP5318	2005-S2
1001	INFS6014	2004-S1
1003	ISYS3207	2005-S2

UnitOfStudy		
<u>uos_code</u>	title	points
COMP5138	Relational DBMS	6
COMP5318	Data Mining	6
INFO6007	IT Project Management	6
SOFT1002	Algorithms	12
ISYS3207	IS Project	4
COMP5702	MIT Research Project	18

Exercise 1: Evaluating a Simple Query

Student			
<u>sid</u>	name	gender	country
1001	Ian	M	AUS
1002	Ha Tshi	F	ROK
1003	Grant	M	AUS
1004	Simon	M	GBR
1005	Jesse	F	CHN
1006	Franziska	F	GER

Enrolled		
<u>sid</u>	<u>uos_code</u>	semester
1001	COMP5138	2005-S2
1002	COMP5702	2005-S2
1003	COMP5138	2005-S2
1006	COMP5318	2005-S2
1001	INFS6014	2004-S1
1003	ISYS3207	2005-S2

UnitOfStudy		
<u>uos_code</u>	title	points
COMP5138	Relational DBMS	6
COMP5318	Data Mining	6
INFO6007	IT Project Management	6
SOFT1002	Algorithms	12
ISYS3207	IS Project	4
COMP5702	MIT Research Project	18

Using the above database instance, give a corresponding relation that gives the titles of all units worth 6 credit points. Think about the steps you take.

title
Relational DBMS
Data Mining
IT Project Management

How does a RDBMS get the answer?


Give the titles of all units worth 6 credit points

Relational Algebra expression:


$$\pi_{\text{title}}(\sigma_{\text{points}=6}(\text{UnitOfStudy}))$$

UnitOfStudy

<u>uos_code</u>	title	points
COMP5138	Relational DBMS	6
COMP5318	Data Mining	6
INFO6007	IT Project Management	6
SOFT1002	Algorithms	12
ISYS3207	IS Project	4
COMP5702	MIT Research Project	18



<u>uos_code</u>	title	points
COMP5138	Relational DBMS	6
COMP5318	Data Mining	6
INFO6007	IT Project Management	6




title
Relational DBMS
Data Mining
IT Project Management




Visualisation of Relational Algebra


Projection (π)




Selection (σ)




Cross-product (\times)




Join (\bowtie)




Set Union (\cup)



Set Minus ($-$)



Set Intersection (\cap)



- › Operators take one/more relations as inputs and give a new relation as result
- › Operations can be chained together to form expressions (queries)
- 1. Operations that remove parts of a relation
 - **Selection** (σ) selects a subset of rows from relation.
 - **Projection** (π) deletes unwanted columns from relation.
- 2. Operations that combine tuples from two relations
 - **Cross-product** (\times) to combine *every* tuple from two relations.
 - **Join** (\bowtie) to combine *matching* tuples from two relations.
- 3. Set Operations
 - **Union** (\cup) tuples in relation 1 or in relation 2.
 - **Intersection** (\cap) tuples in relation 1, as well as in relation 2.
 - **Difference** ($-$) tuples in relation 1, but not in relation 2.
- 4. A schema-level 'rename' operation
 - **Rename** (ρ) allows us to rename a field or relation.

- › ‘Extracts’ columns for attributes that are in *projection* list.
 - Schema of result contains exactly the fields in the projection list, with the same names that they had in the input relation.
- › Examples:

$\pi_{name, country} (Student)$

Student	
name	country
Ian	AUS
Ha Tsch	ROK
Grant	AUS
Simon	GBR
Jesse	CHN
Franziska	GER

$\pi_{title} (UnitOfStudy)$

UnitOfStudy
title
Relational DBMS
Data Mining
IT Project Management
Algorithms
IS Project
MIT Research Project

- › Selects rows that satisfy a ***selection condition***.

- Example:

$$\sigma_{\text{country}='AUS'}(\text{Student})$$

<i>Student</i>			
<u>sid</u>	name	gender	country
1001	Ian	M	AUS
1003	Grant	M	AUS

- › Result relation can be the input for another relational algebra operation! (***Operator composition***.)

- Example:

$$\Pi_{\text{name}}(\sigma_{\text{country}='AUS'}(\text{Student}))$$

<i>Student</i>
name
Ian
Grant

- › Defined as: $R \times S = \{ts \mid t \in R \wedge s \in S\}$
 - each tuple of R is paired with each tuple of S .
 - Resulting schema has one field per field of R and S , with field names 'inherited' if possible.
 - It might end in a conflict with two fields of the same name -> rename needed
- › Sometimes also called *Cartesian product*

› Example:

R			S				result				
A	B		C	D	E		A	B	C	D	E
α	1	X	α	10	a	=	α	1	α	10	a
β	2		β	10	a		α	1	β	10	a
			β	20	b		α	1	β	20	b
			γ	10	b		α	1	γ	10	b
							β	2	α	10	a
							β	2	β	10	a
							β	2	β	20	b
							β	2	γ	10	b

> Conditional Join:

- Example: $R \bowtie_{\theta} S = \sigma_{\theta}(R \times S)$

Student ⋈

Lecturer

Student.f_name = Lecturer.last_name \wedge Student.sid < Lecturer.empid

sid	given	f_name	gender	country	empid	l_name	last_name	room
1001	Cho	Chung	M	AUS	47112344	Vera	Chung	321
1004	Ciao	Poon	M	CHN	12345678	Simon	Poon	431
1004	Ciao	Poon	M	CHN	99004400	Josiah	Poon	482
1111	Alice	Poon	F	AUS	12345678	Simon	Poon	431
1111	Alice	Poon	F	AUS	99004400	Josiah	Poon	482
...	

- > Result schema same as the cross-product's result schema.
- > Sometimes called *theta-join*.
- > **Equi-Join**: Special case where the condition θ contains only equalities.

- › Can compose multiple expressions using logical connectives:
 - \wedge - means AND
 - \vee - means OR
 - \neg - means NOT
- › You can also include comparisons between constants and columns of the input to an operator:
 - Can use $<$, $>$, \leq , \geq , \neq , $=$

› Natural Join: $R \bowtie S$

- Equijoin on all common fields.
- Result schema similar to cross-product, but only one copy of fields for which equality is specified. (according to Ramakrishnan & Gherke also holds for equijoin, but some other texts don't require this for pure equijoin)

<i>Enrolled</i>	
<u>sid</u>	<u>uos_code</u>
1001	COMP5138
1002	COMP5702
1003	COMP5138
1006	COMP5318
1001	INFO6007
1003	ISYS3207



<i>UnitOfStudy</i>		
<u>uos_code</u>	title	points
COMP5138	Relational DBMS	6
COMP5318	Data Mining	6
INFO6007	IT Project Mgmt.	6
SOFT1002	Algorithms	12
ISYS3207	IS Project	4
COMP5702	MIT Research Project	18



<i>result</i>			
sid	uos_code	title	points
1001	COMP5138	Relational DBMS	6
1002	COMP5702	MIT Research Project	18
1003	COMP5138	Relational DBMS	6
1006	COMP5318	Data Mining	6
1001	INFO6007	IT Project Mgmt.	6
1003	ISYS3207	IS Project	4


- › These operations take two input relations R and S
 - Set Union $R \cup S$
 - Definition: $R \cup S = \{ t \mid t \in R \vee t \in S \}$
 - Set Intersection $R \cap S$
 - Definition: $R \cap S = \{ t \mid t \in R \wedge t \in S \}$
 - Set Difference $R - S$
 - Definition: $R - S = \{ t \mid t \in R \wedge t \notin S \}$
- › Important constraint: R and S have the same schema
 - R, S have the *same arity* (same number of fields)
 - ‘Corresponding’ fields must have the same names and domains

Exercise 2: Set Operations

- Suppose you have the following relation. Use a set operation in an RA expression to return all students who are not postgraduates.

$\pi_{\text{sid}}(\text{Student}) - \text{Postgraduate}$

Postgraduate
<u>sid</u>
1003
1004
1005



<u>sid</u>
1001
1002
1006

- › Allows us to name, and therefore to refer to, the results of relational-algebra expressions.
- › Allows us to refer to a relation by more than one name.
- › Notation 1: $\rho_X(E)$
 - returns the expression E under the name X
- › Notation 2: $\rho_X(A1, A2, \dots, An)(E)$
 - returns the result of expression E under the name X , and with the attributes renamed to $A1, A2, \dots, An$.
 - (assumes that the relational-algebra expression E has arity n)
- › Example: $\rho_{\text{Classlist}(student, cid, title, credits)}(Enrolled \times UnitOfStudy)$

Basic Versus Derived Operations

- › We can distinguish between basic and derived RA operators
- › Only 6 basic operators are required to express everything else:
 - **Union** (\cup) tuples in relation 1 or in relation 2.
 - **Set Difference** ($-$) tuples in relation 1, but not in relation 2.
 - **Selection** (σ) selects a subset of rows from relation.
 - **Projection** (π) deletes unwanted columns from relation.
 - **Cross-product** (\times) allows us to fully combine two relations.
 - **Rename** (ρ) allows us to rename one field to another name.
- › Additional (derived) operations:
 - Eg: intersection, join:
 - Not essential, but (very!) useful.
 - E.g., Join: $R \bowtie_{\theta} S = \sigma_{\theta} (R \times S)$

- › Ramakrishnan/Gehrke (3rd edition - the 'Cow' book (2003))
 - Chapter 4 (you can skip the sections on Relational Calculus)
one compact section on RA, including a discussion of relational division
- › Kifer/Bernstein/Lewis (2nd edition – 2006)
 - Chapter 5.1
one section on RA that covers everything as discussed here in the lecture
- › Ullman/Widom (3rd edition – 2008)
 - Chapter 2.4
a nice and gentle introduction to the basic RA operations, leaves out relational division though
 - Chapters 5.1 and 5.2
goes beyond what we cover here in the lecture by extending RA to bags and also introduces grouping, aggregation and sorting operators

- › Foundations of Declarative Querying
 - Relational Algebra
 - Six basic operations
 - Join operation
 - Composition and equivalence rules
- › **Introduction to SQL**
 - Overview
 - Basic SQL Queries
 - Joins Queries
 - Aggregate Functions and Set Operations

SQL: The Structured Query Language

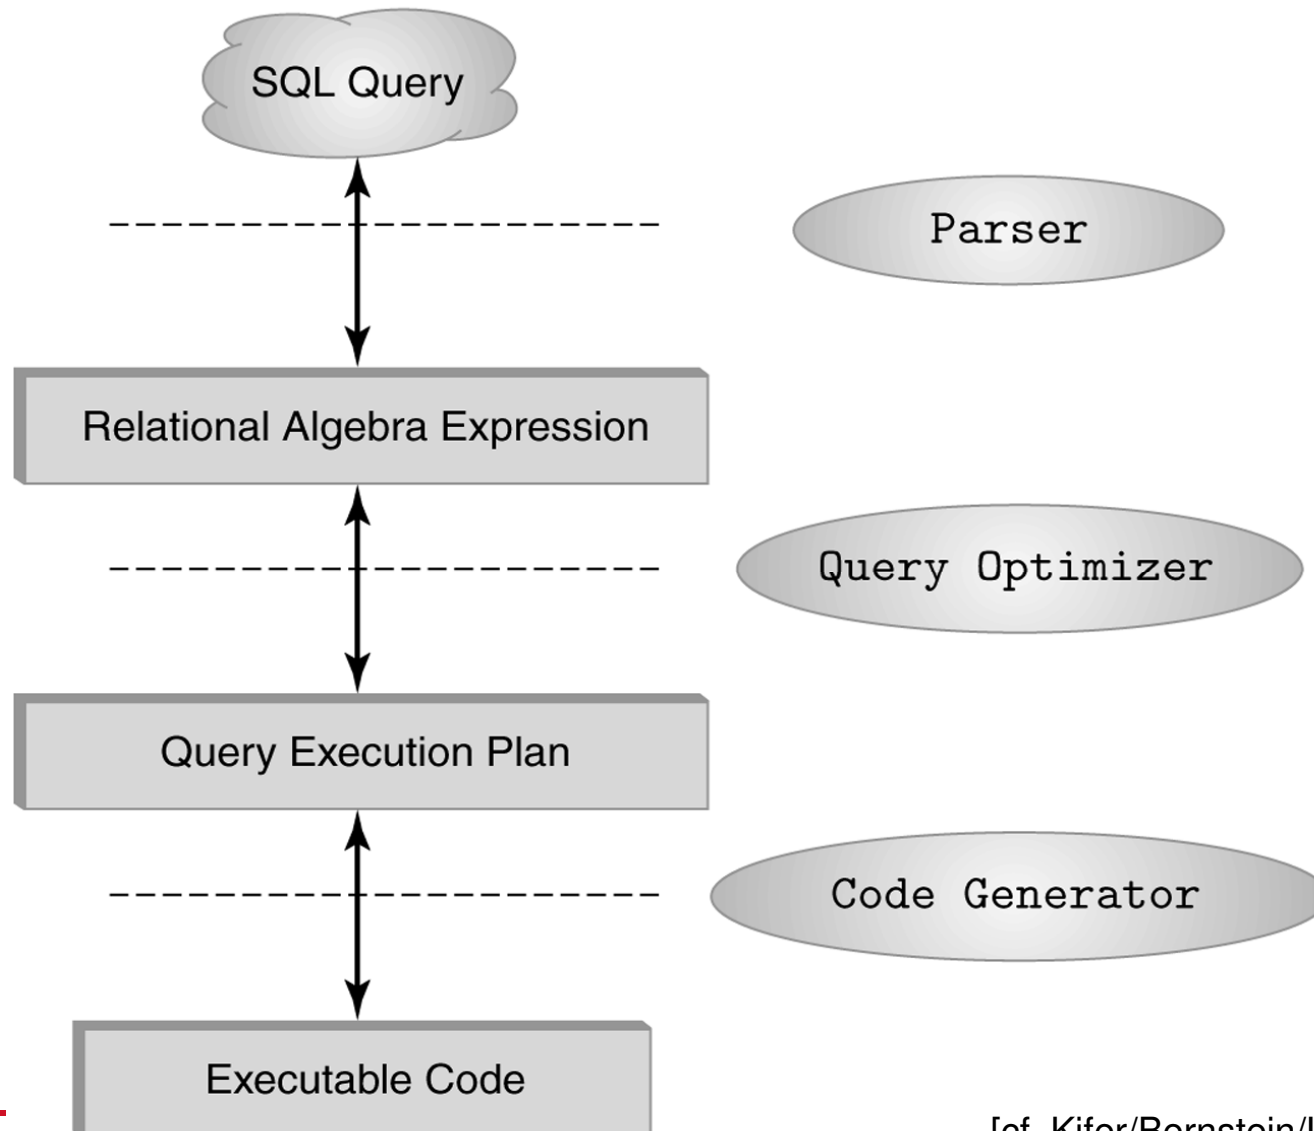
- › SQL is the standard *declarative* query language for RDBMS
 - Describing *what* data we are interested in, but *not how* to retrieve it.
- › Based on SEQUEL
 - Introduced in the mid-1970's as the query language for IBM's System (Structured English Query Language)
- › ANSI standard since 1986, ISO-standard since 1987
- › 1989: revised to SQL-89
- › 1992: more features added – **SQL-92**
- › 1999: major rework – **SQL:1999** (SQL 3)
- › SQL:2003 – 'bugfix release' of SQL:1999 plus SQL/XML
- › SQL:2008 – slight improvements, e.g. INSTEAD OF triggers
- › SQL:2011 – quite a few new time based features, eg. time-period definitions

- › **DDL** (Data Definition Language)
 - Create, drop, or alter the relation schema
- › **DML** (Data Manipulation Language)
 - The retrieval of information stored in the database
 - A **Query** is a statement requesting the retrieval of information
 - The portion of a DML that involves information retrieval is called a **query language**
 - The insertion of new information into the database
 - The deletion of information from the database
 - The modification of information stored in the database
- › **DCL** (Data Control Language)
 - Commands that control a database, including administering privileges and users

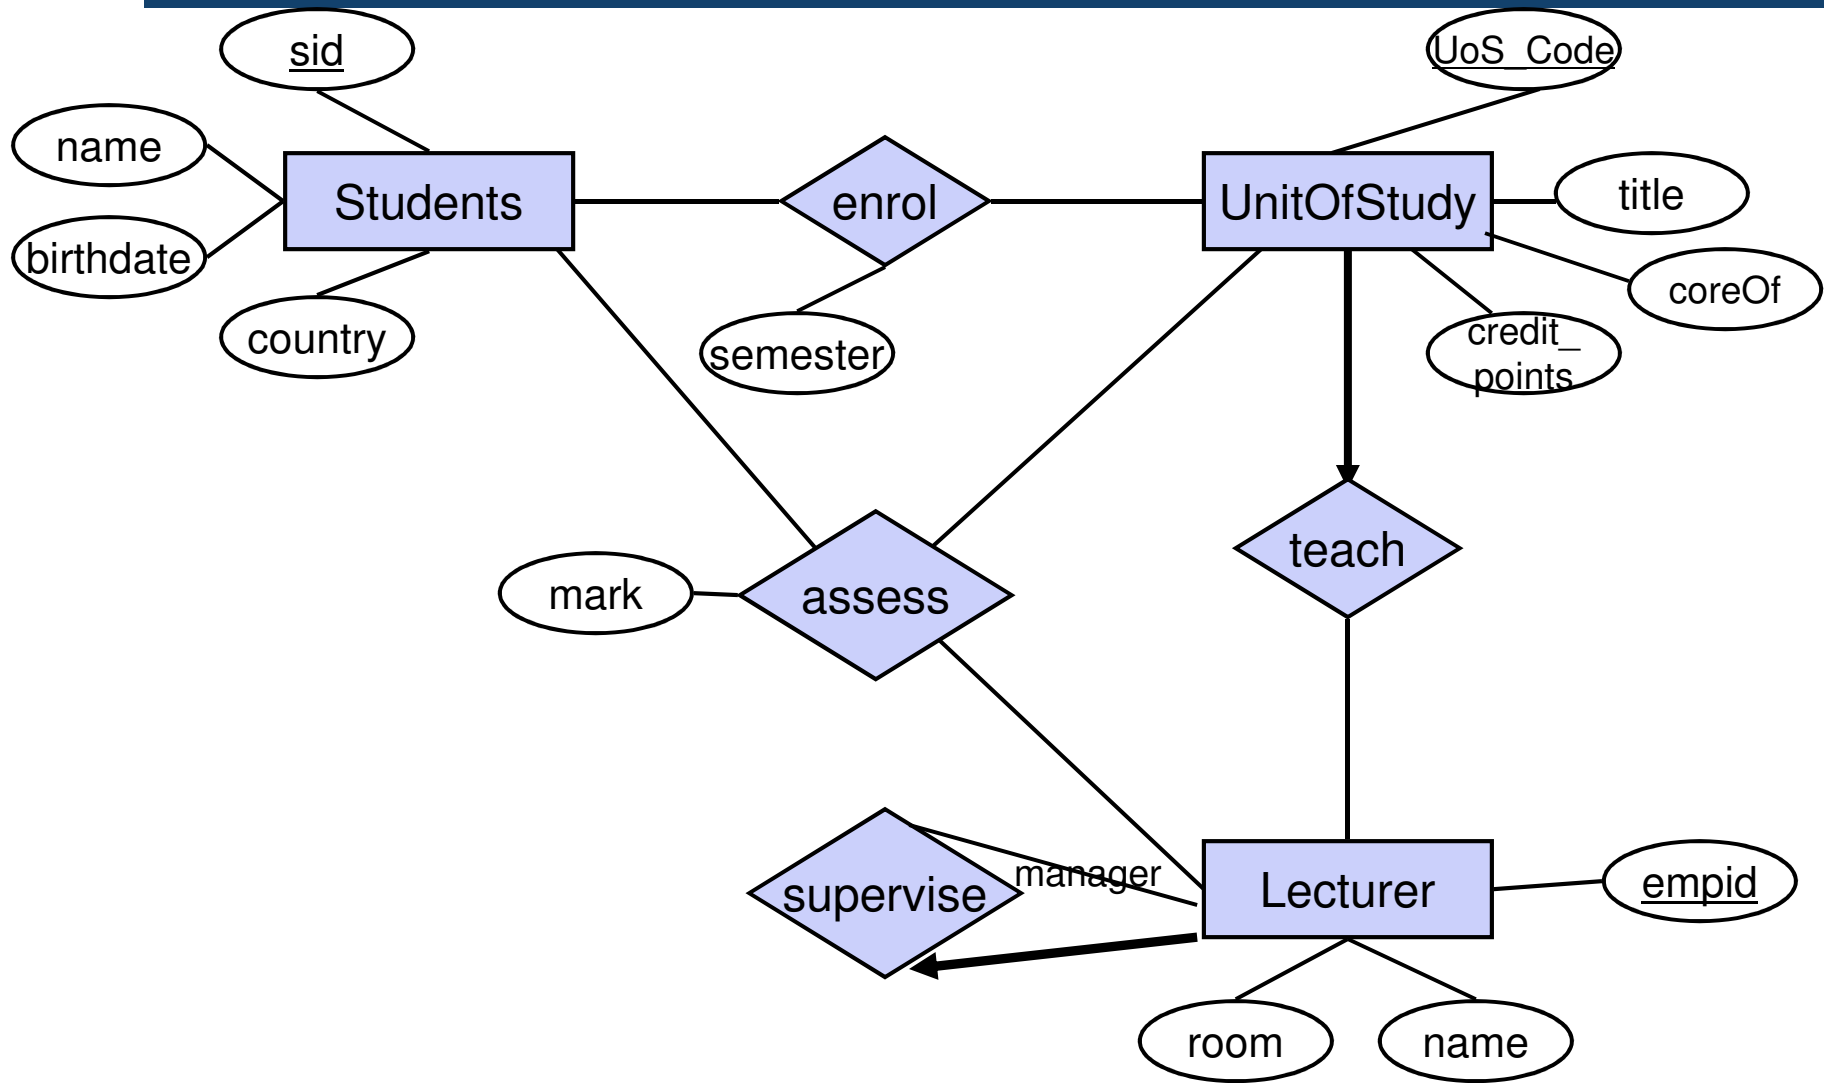
- › Used for queries on single or multiple tables
- › Clauses of the SELECT statement:
 - **SELECT** Lists the columns (and expressions) that should be returned from the query
 - **FROM** Indicate the table(s) from which data will be obtained
 - **WHERE** Indicate the conditions to include a tuple in the result
 - **GROUP BY** Indicate the categorization of tuples
 - **HAVING** Indicate the conditions to include a category
 - **ORDER BY** Sorts the result according to specified criteria
- › The result of an SQL query is a relation
- › a Select-From-Where (SFW) query is equivalent to the relational algebra expression:
$$\pi_{A_1, A_2, \dots, A_n} (\sigma_{condition} (R_1 \times R_2 \times \dots \times R_m))$$



The Role of SQL & RA in RDBMS

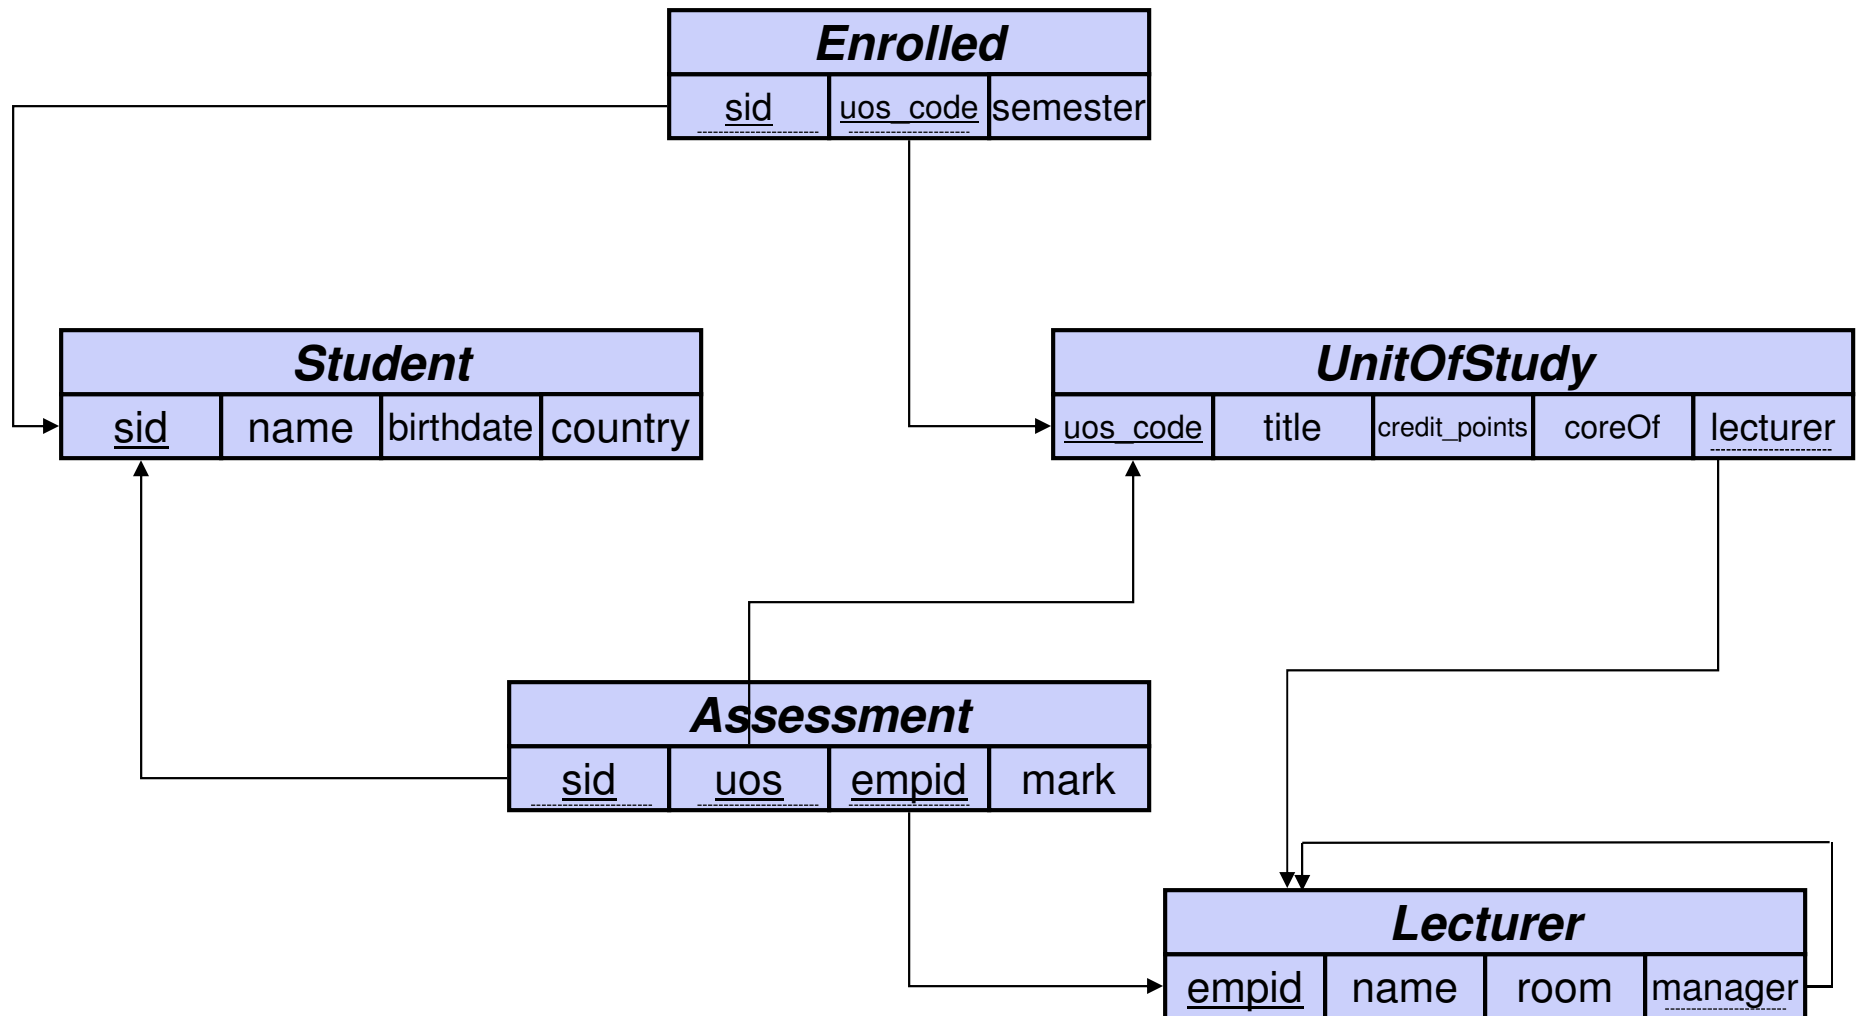


Running Example: ER Diagram





Running Example: Database Schema



Example: Basic SQL Query

- › List the names of all Australian students.

```
SELECT name FROM Student WHERE country='AUS'
```

- › 'Corresponding' relational algebra expression

$$\pi_{name} (\sigma_{country='AUS'} (Student))$$

- › Note: RDBMS' can return duplicate rows
 - (eg: having same values in the name column above)
 - Sets containing duplicate entries are sometimes called multi-sets
 - Strict Relational Algebra will only return sets (no duplicate values)

Example: Order By Clause

- › List all students (name) from Australia in alphabetical order.

```
select name  
from Student  
where country='AUS'  
order by name
```

- › Two options (per attribute):
 - **ASC** ascending order (default)
 - **DESC** descending order
- › You can order by more than one attribute
 - e.g., **order by** country **desc**, name **asc**

- › In contrast to the relational algebra, SQL allows duplicates in relations as well as in query results.
- › To force the elimination of duplicates, insert the keyword **distinct** after **select**.
- › Example: List the countries where students come from.

```
select distinct country  
from Student
```

- › The keyword **all** specifies that duplicates not be removed.

```
select all country  
from Student
```

Arithmetic Expressions in Select Clause

- › An asterisk in the select clause denotes „all attributes“

```
SELECT *  
FROM Student
```

- › The select clause can obtain arithmetic expressions involving the operations +, -, * and /, and operating on constants or attributes of tuples.
- › The query:

```
SELECT uos_code, title, credit_points*2, coreOf, lecturer  
FROM UnitOfStudy
```

would return a relation which is the same as the *UnitofStudy* relation except that the credit-point-values are doubled.

- › SQL allows renaming relation attributes using the **as** clause:

old_name as new_name

- › This is very useful to give, e.g., result columns of expressions a meaningful name.
- › can also assign names to relations as shown in example below
- › Example:
 - Find the student id, employee id and lecturer of all assessments for PHYS1001; rename the column name *empid* as *lecturer*.

```
select a.sid, a.empid as lecturer, a.mark  
from Assessment a  
where a.uos = 'PHYS1001'
```


- › The where clause specifies conditions that the result must satisfy
 - ‘corresponds’ to the selection predicate of the relational algebra.
- › Comparison operators in SQL: = , > , >= , < , <= , <> (or !=)
- › Comparison results can be combined using the logical connectives **AND**, **OR**, and **NOT**.
- › Comparisons can be applied to results of arithmetic expressions
- › Example: Find all UoS codes for units taught by lecturer 1011 that are worth more than four credit points:

```
SELECT uos_code
FROM UnitOfStudy
WHERE lecturer = 1011 AND credit_points > 4
```

- › The **from** clause lists the relations involved in the query
 - corresponds to the Cartesian product operation of the relational algebra.
 - join-predicates explicitly stated in the **where** clause

- › Examples:

- Find the Cartesian product *Student* x *UnitOfStudy*

```
SELECT *  
FROM Student, UnitofStudy
```

- Find the student ID, name, and gender of all students enrolled in INFO2120:

```
SELECT sid, name, gender  
FROM Student, Enrolled  
WHERE Student.sid = Enrolled.sid AND  
      uos_code = 'INFO2120'
```

- › Which students did enroll in what semester?

Join involves multiple tables in
FROM clause

```
SELECT S.sid, S.name, E.semester  
FROM Student S, Enrolled E  
WHERE S.sid = E.sid;
```

WHERE clause performs the
equality check for common
columns of the two tables

- › Some queries need to refer to the same relation twice
- › In this case, aliases are given to the relation name
 - Example: For each academic, retrieve the academic's name, and the name of his or her immediate supervisor.

```
SELECT      L.name, M.name
FROM        Lecturer L, Lecturer M
WHERE       L.manager = M.empid
```

- We can think of **L** and **M** as two different copies of **Lecturer**; **L** represents lecturers in role of supervisees and **M** represents lecturers in role of supervisors (managers)

- › Overview
- › Basic SQL Queries
- › **Join Queries**
- › Aggregate Functions and Set Operations



- › **Join** – a relational operation that causes two or more tables to be combined into a single table
- › **Theta join** – a join in which tuples are combined if they meet some condition (denoted θ)
- › **Equi-join** – a simple case of a theta join in which the joining condition is based on equality between values in the common columns (columns with same name/domain) – common columns appear only once (Ramakrishnan & Gherke)
- › **Natural join** – an equi-join in which we consider all common columns (columns with same name/domain) – common columns appear only once
- › **Left / Right Outer join** – *a join in which rows that do not have matching values in common columns are nonetheless also included (exactly once) in the result table (either rows from the relation to the left or right of the join expression that do not have matching values in common columns are also included – depending on whether we have a left or right outer join.)*
 - *(as opposed to inner join, in which rows must have matching values in order to appear in the result table)*
- › **Full outer join** – *includes rows that would be found in an equi-join as well as rows from the left relation that don't have matches, and rows from the right relation that don't have matches*

The common columns in joined tables are usually the primary key of the dominant table and the foreign key of the dependent table in 1:M relationships

- SQL offers join operators to directly formulate the natural join, equi-join, and the theta join RA operations.
- R **natural join** S
- R **inner join** S **on** <join condition>
- R **inner join** S **using** (<list of attributes>)
- These additional operations are typically used as subquery expressions in the from clause

- List all students and in which courses they enrolled.

```
select name, uos_code, semester  
from Student natural join Enrolled
```

- Who is teaching "INFO2120"?

```
select name  
from UnitOfStudy inner join Lecturer on lecturer=empid  
where uos_code='INFO2120'
```

› Available join types:

- inner join
- left outer join
- right outer join
- full outer join

› Join Conditions:

- **natural**
- **on** <join condition>
- **using** <attribute list>

e.g: *Student* **inner join** *Enrolled* **using** (*sid*)

inner join result					
<u>sid</u>	name	birthdate	country	uos codes	semester
112	'A'	01.01.84	India	SOFT1	1
200	'B'	31.5.79	China	COMP2	2

e.g : *Student* **left outer join** *Enrolled* **using** (*sid*)

left outer join result					
<u>sid</u>	name	birthdate	country	uos codes	semester
112	'A'	01.01.84	India	SOFT1	1
200	'B'	31.5.79	China	COMP2	2
210	'C'	29.02.82	Australia	null	null

- › Overview
- › Basic SQL Queries
- › Join Queries
- › **Aggregate Functions and Set Operations**



- › These functions operate on the *multiset* of values of a column of a relation, and return a value

avg: average value

min: minimum value

max: maximum value

sum: sum of values

count: number of values

- › Must use **distinct** in addition to aggregate over *sets*

Note: with aggregate functions you can't have single-valued columns included in the **select** clause

Examples of Aggregate Operators/Functions

- › How many students enrolled?

```
select count(*) from Enrolled  
select count(distinct sid) from Enrolled
```

- › Which was the best mark for 'SOFT2007'?

```
select max(mark)  
from Assessment  
where uos_code = 'SOFT2007'
```

- › What was the average mark for SOFT2007?

```
select avg(mark)  
from Assessment where uos_code='SOFT2007'
```

- › It is possible for tuples to have a null value, denoted by **null**, for some of their attributes
 - Integral part of SQL to handle missing / unknown information
 - **null** signifies that a value *does not exist*, it does *not mean* “0” or “blank”!
- › The predicate **is null** can be used to check for null values
 - e.g. Find students who don't have a mark for an assessment yet.

```
SELECT sid
FROM Assessment
WHERE mark IS NULL
```
- › Consequence: Three-valued logic
- › The result of any arithmetic expression involving null is null
 - e.g. 5 + null returns null
- › However, (most) aggregate functions simply ignore nulls

NULL Values and Three Valued Logic

- › Any comparison with *null* returns *unknown*
 - e.g. $5 < null$ or $null <> null$ or $null = null$
- › Three-valued logic using the truth value *unknown*:
 - OR: (*unknown* **or** *true*) = *true*, (*unknown* **or** *false*) = *unknown*
(*unknown* **or** *unknown*) = *unknown*
 - AND: (*true* **and** *unknown*) = *unknown*, (*false* **and** *unknown*) = *false*, (*unknown* **and** *unknown*) = *unknown*
 - NOT: (**not** *unknown*) = *unknown*
- › Result of **where** clause predicate is treated as false if it evaluates to unknown
 - e.g: **select** sid **from** assessment **where** mark < 50
ignores all students without a mark so far

- › Aggregate functions except **count(*)** ignore null values on the aggregated attributes
 - result is null if there is no non-null amount

- › Examples:

- Average mark of all assignments

```
SELECT AVG (mark)
FROM Assessment
```

-- ignores tuples with nulls

- Number of all assignments

```
SELECT COUNT (*)
FROM Assessment
```

-- counts *all* tuples (only with *)

- › The set operations **union**, **intersect**, and **except** (Oracle: **minus**) operate on relations and correspond to the relational algebra operations \cup , \cap , $-$.
- › Each of the above operations automatically eliminates duplicates; to retain all duplicates use the corresponding multiset versions **union all**, **intersect all** and **except all**.

Suppose a tuple occurs m times in r and n times in s , then, it occurs:

- $m + n$ times in r **union all** s
- $\min(m, n)$ times in r **intersect all** s (*)
- $\max(0, m - n)$ times in r **except all** s (*)

(*) not supported by Oracle

- › Find all customer names that have a loan, an account, or both:

```
(select customer_name from depositor)  
  union  
(select customer_name from borrower)
```

- › Find all customer names that have both a loan and an account

```
(select customer_name from depositor)  
  intersect  
(select customer_name from borrower)
```

- › Find all customer names that have an account but no loan

```
(select customer_name from depositor)  
  except  
(select customer_name from borrower)
```


Examples for Set Operations

- › List all names in the database.

```
select name from Student
union
select name from Lecturer
```

- › Which students did not enroll in any units?

```
select sid from Student
minus
select sid from Enrolled
```

- › Find students who enrolled for 'COMP5138' and 'COMP5318'.

```
select sid from Enrolled where uoscode='COMP5138'
intersect
select sid from Enrolled where uoscode='COMP5318'
```

› Foundations of SQL

- Relational Algebra:
Projection, Selection, Rename, Set Operations, Cross Product,
Joins (equi-join, theta-join, natural join)
- how it relates to SQL queries

› Introduction to SQL

- SELECT ... FROM ... WHERE ... ORDER BY ...
- **Joins** in SQL
- **NULL values and semantic**

› Aggregate Functions and Set Operations

- Count, Sum, Min, Max, Avg, ...
- Union, Intersect, and Except

Each database textbook has a pretty standard chapter on SQL that covers all commands that we discussed in this lecture:

- › Ramakrishnan/Gehrke (3rd edition - the 'Cow' book (2003))
 - Chapter 5
 - uses the famous 'Sailor-database' as examples*
- › Kifer/Bernstein/Lewis (2nd edition – 2006)
 - Chapter 5
 - includes some helpful visualisations on how complex SQL is evaluated*
- › Ullman/Widom (3rd edition – 2008)
 - Chapter 6
 - up-to 6.5 good introduction and overview of all parts of SQL querying*
- › Silberschatz/Korth/Sudarshan (5th edition - 'sailing boat')
 - Sections 3.1-3.6
- › Elmasri/Navathe (5th edition)
 - Sections 8.4 and 8.5.1

- › Integrity Constraints
 - Domain and CHECK constraints
 - ON DELETE and ON UPDATE actions, deferred constraints
 - Assertions
 - Triggers
- › Readings :
 - Ramakrishnan/Gehrke (3rd edition - the 'Cow' book)
 - Sections 3.2-3.3 and Sections 5.7-5.9
 - *Integrity constraints are covered in different parts of the SQL discussion; only brief on triggers*
 - Kifer/Bernstein/Lewis (2nd edition)
 - Sections 3.2.2-3.3 and Chapter 7
 - *Integrity constraints are covered as part of the relational model, but a good dedicated chapter (Chap 7) on triggers*
 - Ullman/Widom (3rd edition)
 - Chapter 7
 - *Has a complete chapter dedicated to both integrity constraints&triggers.*

- › SQL includes a between comparison operator (called “range queries”)
- Example: Find all students (by SID) who gained marks in the distinction and high-distinction range in COMP5138.

```
SELECT sid
  FROM Assessment
 WHERE uos_code = 'COMP5138' AND
       mark BETWEEN 75 AND 100
```

- › Find the students with marks between 0 and 10.

```
select sid  
from Assessment  
where mark between 0 and 10
```

- › Who is teaching “INFO2120”?

```
select name  
from UnitsOfStudy, Lecturer  
where uos_code='INFO2120' and lecturer=empid
```

- › List students who are enrolled in INFO2120 in either semester 1 or 2

```
select sid from Enrolled  
where uos_code ='INFO2120' and semester in (1, 2)
```

- › SQL includes a string-matching operator for comparisons on character strings.
 - **LIKE** is used for string matching
- › Patterns are described using two special characters (“wildcards”):
 - percent (%). The % character matches any substring.
 - underscore (_). The _ character matches any character.
- › List the titles of all “COMP” unit of studies.

```
select title
  from UnitOfStudy
 where uos_code like 'COMP%'
```

- › SQL supports a variety of string operations such as
 - concatenation (using “||”)
 - converting from upper to lower case (and vice versa)
 - finding string length, extracting substrings, etc.

Regular Expression Matches

- › New since SQL:2003: regular expression string matching
 - typically implemented as set of SQL functions, e.g. **regexp_like(...)**
- › What are regular expressions?
 - Pattern consisting of *character literals* and/or *metacharacters*
 - *metacharacters* specify how to process a regular expression
 - () grouping
 - | alternative
 - [] character list
 - . matches any character
 - * repeat preceeding pattern zero, one, or more times
 - + repeat preceeding pattern one or more times
 - ^ start of a line
 - \$ end of line

- › Example:

```
select title
from UnitOfStudy
where regexp_like(uos_code, '^COMP[:digit:]{4}')
```


SQL Type	Example	Accuracy	Description
DATE	'2012-03-26'	1 day	a date (some systems incl. time)
TIME	'16:12:05'	ca. 1 ms	a time, often down to nanoseconds
TIMESTAMP	'2012-03-26 16:12:05'	ca. 1 sec	Time at a certain date: SQL Server: DATETIME
INTERVAL	'5 DAY'	years - ms	a time duration

› Comparisons

- Normal time-order comparisons with '=', '>', '<', '<=', '>=', ...

› Constants

- CURRENT_DATE db system's current date
- CURRENT_TIME db system's current timestamp

› Example: Which students enrolled before today?

- › Database systems support a variety of date/time related ops
 - Unfortunately not very standardized – a lot of slight differences
- › Main Operations
 - **EXTRACT**(*component* **FROM** *date*)
 - e.g. EXTRACT(year FROM enrolmentDate)
 - **DATE** *string* (Oracle syntax: TO_DATE(*string*,*template*))
 - e.g. DATE '2012-03-01'
 - Some systems allow templates on how to interpret *string*
 - Oracle syntax: TO_DATE('01-03-2012', 'DD-Mon-YYYY')
 - **+/- INTERVAL**:
 - e.g. '2012-04-01' + INTERVAL '36 HOUR'
- › Many more -> check database system's manual

**TO BE
CONTINUED...** 